

Comparison of Dimensionality Reduction Techniques on Machine Learning

By: Syed Bilal Rizwan

1. Aim

The aim of this project is to compare different dimensionality reduction techniques and their effect on Machine Learning performance. The techniques are tried on 15 datasets to see the general behavior.

2. Datasets Chosen

2.1 Classification Datasets Description:

1. **Marketing Campaign Dataset:** A response model can provide a significant boost to the efficiency of a marketing campaign by increasing responses or reducing expenses. The objective is to predict who will respond to an offer for a product or service.
2. **Credit Card Fraud Dataset:** The dataset contains transactions made by credit cards in September 2013 by European cardholders. This dataset presents transactions that occurred in two days, where we have 492 frauds out of 284,807 transactions. The dataset is highly unbalanced, the positive class (frauds) account for 0.172% of all transactions.
3. **Heart Disease Prediction Dataset:** According to the CDC, heart disease is one of the leading causes of death for people of most races in the US (African Americans, American Indians and Alaska Natives, and white people). Originally, the dataset come from the CDC and is a major part of the Behavioral Risk Factor Surveillance System (BRFSS), which conducts annual telephone surveys to gather data on the health status of U.S. residents. It consists of 401,958 rows and 279 columns which are reduced to 20 columns.
4. **Diabetes Dataset:** The dataset represents 10 years (1999-2008) of clinical care at 130 US hospitals and integrated delivery networks. It includes over 50 features representing patient and hospital outcomes.
5. **High Income Prediction:** Extraction was done by Barry Becker from the 1994 Census database. A set of clean records was extracted. Prediction task is to determine whether a person makes over 50K a year.
6. **Dry Beans Dataset:** Images of 13,611 grains of 7 different registered dry beans were taken with a high-resolution camera. A total of 16 features; 12 dimensions and 4 shape forms, were obtained from the grains. Prediction task is to find out the type of bean it is.
7. **Banknote Authentication Dataset:** Data were extracted from images that were taken from genuine and forged banknote-like specimens. For digitization, an industrial camera usually used for print inspection was used. The final images have 400x 400 pixels.
8. **Audit Data:** The goal of the research is to help the auditors by building a classification model that can predict the fraudulent firm on the basis the present and historical risk factors.

2.2 Regression Dataset Description:

1. **Combined Cycle Power Plant Dataset:** The dataset contains 9568 data points collected from a Combined Cycle Power Plant over 6 years (2006-2011), when the power plant was set to work with full load. Features consist of hourly average ambient variables Temperature (T), Ambient Pressure (AP), Relative Humidity (RH) and Exhaust Vacuum (V) to predict the net hourly electrical energy output (EP) of the plant. Prediction task is to predict the Electrical Energy Output.
2. **Energy Efficiency Dataset:** Perform energy analysis using 12 different building shapes simulated in Ecotect. The buildings differ with respect to the glazing area, the glazing area distribution, and the orientation, amongst other parameters.
3. **QSAR Aquatic Toxicity Dataset:** This dataset was used to develop quantitative regression QSAR models to predict acute aquatic toxicity towards the fish *Pimephales promelas* (fathead minnow) on a set of 908 chemicals. to predict acute aquatic toxicity towards *Daphnia Magna*. LC50 data, which is the concentration that causes death in 50% of test D. magna over a test duration of 48 hours, was used as model response.
4. **Bike Sharing Dataset:** Bike sharing systems are new generation of traditional bike rentals where entire process from membership, rental and return has become automatic. Through these systems, user can easily rent a bike from a particular position and return at another position. Currently, there are about over 500 bike-sharing programs around the world which is composed of over 500 thousand bicycles. Today, there exists great interest in these systems due to their key role in traffic, environmental and health issues. Goal is to predict the number of bikes given the other variables.
5. **Wine Quality Dataset:** Goal is to predict the quality of wine given the other variables.

6. **Student Performance Dataset:** This data approach student achievement in secondary education of two Portuguese schools. The data attributes include student grades, demographic, social and school related features) and it was collected by using school reports and questionnaires.
7. **Buzz in social media(Tom's Hardware Dataset):** This dataset contains examples of buzz events from two different social networks: Twitter, and Tom's Hardware, a forum network focusing on modern technology with more conservative dynamics.

3. Background

Training a machine learning model on large datasets require a lot of computational resources and is excessively time consuming as well. To achieve the end-goal in a realistic timeframe, it is important to think of ways to pre-process dataset in a way which leads to less computation and allow scalability. This is where dimensionality reduction techniques come into the picture.

Dimensionality reduction maps a high dimensional dataset into a lower dimensional space without losing information in the dataset. These techniques are used as a pre-processing step before using the dataset for training. There are several renowned DR techniques of which a few are chosen for the comparison analysis below.

3.1 Principal Component Analysis (PCA):

3 different variants of PCA are tried below.

3.1.1 Normal PCA:

PCA converts high dimensional dataset into a lower dimensional dataset while still capturing maximum information in the dataset. It does that by converting n correlated features into k uncorrelated features(components) where k is smaller than n . Furthermore, it ensures that maximum variance is captured by the first component and the second highest variance is captured by the second principal component so this way most of the variance is captured by the first few independent components which are then used to transform the dataset into a lower dimension.

3.1.2 Sparse PCA(sPCA):

Sparse PCA is another variant of PCA that extract sparse components which can help in reconstruction of data. It overcomes the disadvantage of normal PCA which uses all input features to generate the transformed data, but sparse PCA only uses a few input features to transform the data.

3.1.3 Incremental PCA(ipca):

This is like normal PCA however incremental PCA is for large datasets which might be too large to fit to the memory. Incremental PCA makes a low rank approximation which is not based on the number of samples but only on the number of features.

3.2 Linear Discriminant Analysis (LDA):

Linear Discriminant Analysis utilizes class labels along with the dataset to reduce dimensionality making it a supervised dimensionality reduction technique as opposed to PCA. It is a technique that is used to find linear combination of features that ensure separability of classes. Furthermore, the number of components found are always less than number of classes which means it is a strong dimensionality reduction technique. For example, if LDA was applied on a binary classification dataset, then the resulting components would just be 1. Lastly, this technique is only applicable on classification datasets.

3.3 Singular value Decomposition(SVD):

This technique is like PCA where the only difference is that the matrix factorization is performed on data matrix rather than the covariance matrix which is the case for PCA.

4. Methodology

Analysis is done on 15 datasets consisting of 8 classification and 7 regression datasets. The project code structure is divided into these two parts Classification and Regression. Each part is also divided further into more sections. In the first section, the datasets are loaded. In the second section, the datasets are pre-processed and lastly in the last section, each dataset goes through the Machine learning analysis using Dimensionality Reduction Techniques.

4.1 Pre-Processing

Once datasets are loaded, each dataset is pre-processed according to its need. A bird's eye view of original dataset is printed before starting its pre-processing.

1. Firstly, the datatypes and missing values in each column is checked. If any missing values exist, they are addressed by either removing them or imputing them.
2. Unnecessary columns are removed
3. Categorical variables are one-hot encoded
4. All numerical columns are scaled using a Min-Max Scaler
5. If any additional pre-processing is required by any dataset, it is done in the last step.
6. Predictors are separated from target variable. Convention is to name predictors dataframe as dataset_df and target variable as dataset_classes

Finally, the predictors dataset is printed to see how it looks. This process is repeated for all datasets belonging to that section i.e., classification or regression.

4.2 Machine Learning with Dimensionality Reduction

After all datasets belonging to that section are pre-processed, machine learning is carried out by trying out various dimensionality reduction techniques.

1. Firstly, a pipeline is developed to do the whole analysis and give us final results dataframe. Then, the pipeline is run on all the datasets of that part. Lazy predict library is used to automate running different machine learning models for the classification task and regression task. The pipeline can be broken into 6 parts:
 - i) Lazy Predict on dataset with original features
 - ii) Applying PCA and then running Lazy Predict on resulting dataset.
 - iii) Applying other PCA variants and then running Lazy Predict on resulting dataset.
 - iv) Applying LDA and then running resulting dataset (only applicable for classification datasets)
 - v) Applying SVD and then running resulting dataset
 - vi) Compiling results from each iteration and output a results dataframe
2. Then, each dataset is passed through the pipeline and its results are exported into an excel sheet.
3. Results are printed on the notebook and a detailed analysis is done for the results of that dataset.

5. Implementation

5.1 Importing Libraries

```
In [1]: #Importing Supporting Libraries
import numpy as np
import pandas as pd
from statistics import mean
import time
from numpy import *
import warnings
warnings.filterwarnings('ignore')
%load_ext autotime

#Importing Pre-processing Libraries
from pandas.api.types import is_numeric_dtype
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split

#Importing Dimensionality Reduction Libraries
from sklearn.decomposition import PCA, IncrementalPCA, KernelPCA, SparsePCA, TruncatedSVD
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis

#Importing Machine Learning Pipeline
from lazypredict.Supervised import LazyRegressor, LazyClassifier
import lazypredict
```

time: 3.75 s (started: 2022-12-30 00:51:41 +05:00)

5.2 Classification Datasets

5.2.1 Loading Datasets

```
In [8]: marketing_df = pd.read_csv('Classification/MarketingDataUCI.csv', sep='\t') #Marketing Campaign Dataset
credit_df = pd.read_csv('Classification/CreditCardUCI.csv') #Credit Card Fraud Dataset
```

```
heart_df = pd.read_csv('Classification/HeartDataUCI.csv') #Heart disease Dataset
diabetic_df = pd.read_csv('Classification/DiabeticDataUCI.csv') #Diabetes Dataset
income_df = pd.read_csv('Classification/IncomeDataUCI.csv') #High Income Prediction
beans_df = pd.read_excel('Classification/DryBeanDataUCI.xlsx') #Drybeans dataset
banknotes_df = pd.read_csv('Classification/BankNoteAuthenticationUCI.txt',
                           header = None, names = ['variance', 'skewness',
                                                    'curtosis', 'entropy', 'Class']) #Banknotes authentication dataset
audit_df = pd.read_csv('Classification/AuditRiskUCI.csv') #Audit Risk Dataset
```

```
#Printing Shape of each dataset
print('Shape of Marketing dataframe is: ', marketing_df.shape)
print('Shape of credit card dataframe is: ', credit_df.shape)
print('Shape of heart disease dataframe is: ', heart_df.shape)
print('Shape of diabetes dataframe is: ', diabetic_df.shape)
print('Shape of Income dataframe is: ', income_df.shape)
print('Shape of Beans dataframe is: ', beans_df.shape)
print('Shape of Bank Notes dataframe is: ', banknotes_df.shape)
print('Shape of audit dataframe is: ', audit_df.shape)
```

```
Shape of Marketing dataframe is: (2240, 29)
Shape of credit card dataframe is: (284807, 31)
Shape of heart disease dataframe is: (319795, 18)
Shape of diabetes dataframe is: (101766, 50)
Shape of Income dataframe is: (68378, 15)
Shape of Beans dataframe is: (13611, 17)
Shape of Bank Notes dataframe is: (1372, 5)
Shape of audit dataframe is: (776, 27)
time: 4.03 s (started: 2022-12-30 03:32:53 +05:00)
```

5.2.2 Pre-Processing Datasets

1. Marketing dataset

```
In [3]: print(marketing_df.isnull().sum()) #Check missing values
marketing_df.head() #Bird's eye view of dataset
```

```
ID          0
Year_Birth  0
Education   0
Marital_Status  0
Income      24
Kidhome     0
Teenhome    0
Dt_Customer  0
Recency     0
MntWines    0
MntFruits   0
MntMeatProducts  0
MntFishProducts  0
MntSweetProducts  0
MntGoldProds  0
NumDealsPurchases  0
NumWebPurchases  0
NumCatalogPurchases  0
NumStorePurchases  0
NumWebVisitsMonth  0
AcceptedCmp3  0
AcceptedCmp4  0
AcceptedCmp5  0
AcceptedCmp1  0
AcceptedCmp2  0
Complain     0
Z_CostContact  0
Z_Revenue    0
Response     0
dtype: int64
```

```
Out[3]:
```

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome	Teenhome	Dt_Customer	Recency	MntWines	...	NumWebVisitsMonth
0	5524	1957	Graduation	Single	58138.00	0	0	04-09-2012	58	635	...	7
1	2174	1954	Graduation	Single	46344.00	1	1	08-03-2014	38	11	...	5
2	4141	1965	Graduation	Together	71613.00	0	0	21-08-2013	26	426	...	4
3	6182	1984	Graduation	Together	26646.00	1	0	10-02-2014	26	11	...	6
4	5324	1981	PhD	Married	58293.00	1	0	19-01-2014	94	173	...	5

5 rows × 29 columns

time: 15 ms (started: 2022-12-30 00:51:49 +05:00)

In [4]:

```
marketing_classes = marketing_df[['Response']]
marketing_df.drop(columns = ['Response', 'ID', 'Dt_Customer'], inplace = True) #Dropping unnecessary columns

#dummy-encoding (One-hot encoding) the categorical variables
marketing_df = pd.get_dummies(marketing_df, drop_first = True)
marketing_df.shape

#Replacing missing values by Nan
imputer = SimpleImputer(missing_values=np.nan)
imputer = imputer.fit(marketing_df)
marketing_df = pd.DataFrame(imputer.transform(marketing_df), columns = (marketing_df.columns)).astype(marketing_df.dtypes)

#Scaling and One hot Encoding
Scaler = MinMaxScaler()
marketing_df = pd.get_dummies(marketing_df)
marketing_df = pd.DataFrame(Scaler.fit_transform(marketing_df), columns = marketing_df.columns)
print('Shape of df now is: ', marketing_df.shape)
marketing_df.head()
```

Shape of df now is: (2240, 35)

Out[4]:

	Year_Birth	Income	Kidhome	Teenhome	Recency	MntWines	MntFruits	MntMeatProducts	MntFishProducts	MntSweetProducts	...	Ed
0	0.62	0.08	0.00	0.00	0.59	0.43	0.44	0.32	0.66	0.33	...	
1	0.59	0.07	0.50	0.50	0.38	0.01	0.01	0.00	0.01	0.00	...	
2	0.70	0.11	0.00	0.00	0.26	0.29	0.25	0.07	0.43	0.08	...	
3	0.88	0.04	0.50	0.00	0.26	0.01	0.02	0.01	0.04	0.01	...	
4	0.85	0.09	0.50	0.00	0.95	0.12	0.22	0.07	0.18	0.10	...	

5 rows × 35 columns

time: 31 ms (started: 2022-12-30 00:51:49 +05:00)

2. Credit Card Dataset

In [11]:

```
print(credit_df.isnull().sum()) #Check missing values
print(credit_df.dtypes) #Check data types
credit_df.head() #Bird's eye view of dataset
```

```
Time      0
V1        0
V2        0
V3        0
V4        0
V5        0
V6        0
V7        0
V8        0
V9        0
V10       0
V11       0
V12       0
V13       0
V14       0
V15       0
V16       0
V17       0
V18       0
V19       0
V20       0
V21       0
V22       0
V23       0
V24       0
V25       0
V26       0
V27       0
V28       0
Amount    0
Class     0
dtype: int64
Time      float64
V1        float64
V2        float64
V3        float64
V4        float64
V5        float64
V6        float64
```

```

V7      float64
V8      float64
V9      float64
V10     float64
V11     float64
V12     float64
V13     float64
V14     float64
V15     float64
V16     float64
V17     float64
V18     float64
V19     float64
V20     float64
V21     float64
V22     float64
V23     float64
V24     float64
V25     float64
V26     float64
V27     float64
V28     float64
Amount  float64
Class    int64
dtype: object

```

```

Out[11]:
   Time  V1  V2  V3  V4  V5  V6  V7  V8  V9  ...  V21  V22  V23  V24  V25  V26  V27  V28  Amount  Class
0  0.00 -1.36 -0.07 2.54 1.38 -0.34 0.46 0.24 0.10 0.36  ... -0.02 0.28 -0.11 0.07 0.13 -0.19 0.13 -0.02 149.62 0
1  0.00 1.19 0.27 0.17 0.45 0.06 -0.08 -0.08 0.09 -0.26  ... -0.23 -0.64 0.10 -0.34 0.17 0.13 -0.01 0.01 2.69 0
2  1.00 -1.36 -1.34 1.77 0.38 -0.50 1.80 0.79 0.25 -1.51  ... 0.25 0.77 0.91 -0.69 -0.33 -0.14 -0.06 -0.06 378.66 0
3  1.00 -0.97 -0.19 1.79 -0.86 -0.01 1.25 0.24 0.38 -1.39  ... -0.11 0.01 -0.19 -1.18 0.65 -0.22 0.06 0.06 123.50 0
4  2.00 -1.16 0.88 1.55 0.40 -0.41 0.10 0.59 -0.27 0.82  ... -0.01 0.80 -0.14 0.14 -0.21 0.50 0.22 0.22 69.99 0

```

5 rows × 31 columns

time: 31 ms (started: 2022-12-30 03:34:58 +05:00)

```

In [6]: credit_df.Class.value_counts() #Checking Class Distribution of dataset

```

```

Out[6]:
0    284315
1      492
Name: Class, dtype: int64
time: 0 ns (started: 2022-12-25 18:33:59 +05:00)

```

```

In [7]: credit_df = credit_df.sample(50000) #Sampling rows from dataset

```

time: 16 ms (started: 2022-12-25 18:33:59 +05:00)

```

In [8]: credit_classes = credit_df[['Class']]
credit_df.drop(columns = ['Time', 'Class'], inplace = True) #Dropping unnecessary columns

#Scaling and One hot Encoding
Scaler = MinMaxScaler()
credit_df = pd.get_dummies(credit_df)
credit_df = pd.DataFrame(Scaler.fit_transform(credit_df), columns = credit_df.columns)
print('Shape of df now is: ', credit_df.shape)
credit_df.head()

```

Shape of df now is: (50000, 29)

```

Out[8]:
   V1  V2  V3  V4  V5  V6  V7  V8  V9  V10  ...  V20  V21  V22  V23  V24  V25  V26  V27  V28  Amount
0  0.97 0.61 0.77 0.34 0.56 0.51 0.54 0.73 0.47 0.51  ... 0.58 0.57 0.37 0.51 0.27 0.48 0.44 0.65 0.42 0.00
1  0.97 0.60 0.75 0.22 0.57 0.58 0.53 0.75 0.43 0.52  ... 0.59 0.58 0.41 0.51 0.56 0.49 0.34 0.65 0.42 0.00
2  0.99 0.60 0.77 0.32 0.55 0.51 0.54 0.74 0.50 0.51  ... 0.58 0.58 0.41 0.52 0.34 0.39 0.35 0.65 0.42 0.00
3  0.97 0.59 0.81 0.31 0.53 0.52 0.52 0.74 0.50 0.51  ... 0.58 0.59 0.49 0.51 0.43 0.47 0.40 0.65 0.43 0.00
4  0.95 0.60 0.82 0.42 0.54 0.54 0.53 0.74 0.51 0.50  ... 0.59 0.59 0.48 0.51 0.45 0.48 0.36 0.65 0.43 0.01

```

5 rows × 29 columns

time: 31 ms (started: 2022-12-25 18:33:59 +05:00)

3. Heart Disease Dataset

```

In [12]: print(heart_df.isnull().sum()) #Check missing values
print(heart_df.dtypes) #Check data types
heart_df.head() #Bird's eye view of dataset

```

```
HeartDisease      0
BMI               0
Smoking           0
AlcoholDrinking   0
Stroke            0
PhysicalHealth     0
MentalHealth      0
DiffWalking       0
Sex               0
AgeCategory       0
Race              0
Diabetic          0
PhysicalActivity   0
GenHealth         0
SleepTime         0
Asthma            0
KidneyDisease     0
SkinCancer        0
dtype: int64
HeartDisease      object
BMI               float64
Smoking           object
AlcoholDrinking   object
Stroke            object
PhysicalHealth     float64
MentalHealth      float64
DiffWalking       object
Sex               object
AgeCategory       object
Race              object
Diabetic          object
PhysicalActivity   object
GenHealth         object
SleepTime         float64
Asthma            object
KidneyDisease     object
SkinCancer        object
dtype: object
```

```
Out[12]:
```

	HeartDisease	BMI	Smoking	AlcoholDrinking	Stroke	PhysicalHealth	MentalHealth	DiffWalking	Sex	AgeCategory	Race	Diabetic
0	No	16.60	Yes	No	No	3.00	30.00	No	Female	55-59	White	Yes
1	No	20.34	No	No	Yes	0.00	0.00	No	Female	80 or older	White	No
2	No	26.58	Yes	No	No	20.00	30.00	No	Male	65-69	White	Yes
3	No	24.21	No	No	No	0.00	0.00	No	Female	75-79	White	No
4	No	23.71	No	No	No	28.00	0.00	Yes	Female	40-44	White	No

time: 109 ms (started: 2022-12-30 03:35:03 +05:00)

```
In [11]: heart_df['HeartDisease'].value_counts()[1]/
          (heart_df['HeartDisease'].value_counts()[0] +
           heart_df['HeartDisease'].value_counts()[1]) #Checking Class Distribution
```

```
Out[11]: 0.08559545959130067
```

time: 32 ms (started: 2022-12-25 18:33:59 +05:00)

```
In [12]: heart_df = heart_df.sample(50000) #Sampling rows from dataset
```

time: 31 ms (started: 2022-12-25 18:33:59 +05:00)

```
In [13]: heart_classes = heart_df[['HeartDisease']]
          heart_df.drop(columns = ['HeartDisease'], inplace = True) #Dropping unnecessary columns

          #Scaling and One hot Encoding
          Scaler = MinMaxScaler()
          heart_df = pd.get_dummies(heart_df)
          heart_df = pd.DataFrame(Scaler.fit_transform(heart_df), columns = heart_df.columns)
          print('Shape of df now is: ', heart_df.shape)
          heart_df.head()
```

Shape of df now is: (50000, 50)

```
Out[13]:
```

	BMI	PhysicalHealth	MentalHealth	SleepTime	Smoking_No	Smoking_Yes	AlcoholDrinking_No	AlcoholDrinking_Yes	Stroke_No	Stroke_Yes
0	0.19	0.17	0.33	0.22	1.00	0.00	1.00	0.00	1.00	0.
1	0.28	0.17	0.23	0.30	1.00	0.00	1.00	0.00	1.00	0.
2	0.23	0.67	0.03	0.30	1.00	0.00	1.00	0.00	1.00	0.

	BMI	PhysicalHealth	MentalHealth	SleepTime	Smoking_No	Smoking_Yes	AlcoholDrinking_No	AlcoholDrinking_Yes	Stroke_No	Stroke_Yes
3	0.29	0.00	0.00	0.30	0.00	1.00	1.00	0.00	1.00	0.
4	0.13	0.00	0.00	0.22	0.00	1.00	1.00	0.00	1.00	0.

5 rows × 50 columns

time: 78 ms (started: 2022-12-25 18:33:59 +05:00)

In [14]:

```
heart_classes.loc[
    ((heart_classes['HeartDisease'] == 'Yes')), 'HeartDisease'] = 1    #Labeling the Yes case

heart_classes.loc[
    ((heart_classes['HeartDisease'] == 'No')), 'HeartDisease'] = 0    #Labeling the No case
```

time: 16 ms (started: 2022-12-25 18:34:01 +05:00)

4. Diabetes Dataset

In [25]:

```
#Changing diag_1, diag_2 and diag_3 to numeric
diabetic_df[['diag_1', 'diag_2', 'diag_3']] = diabetic_df[['diag_1', 'diag_2', 'diag_3']].apply(pd.to_numeric, errors='coerce')

#Replacing all missing values with Nan
diabetic_df = diabetic_df.replace('?', np.nan)

print(diabetic_df.isnull().sum()) #Check missing values
print(diabetic_df.dtypes)        #Check data types
diabetic_df.head()               #Bird's eye view of dataset
```

```
encounter_id      0
patient_nbr       0
race              2273
gender            0
age              0
weight           98569
admission_type_id 0
discharge_disposition_id 0
admission_source_id 0
time_in_hospital  0
payer_code        40256
medical_specialty 49949
num_lab_procedures 0
num_procedures    0
num_medications   0
number_outpatient 0
number_emergency  0
number_inpatient  0
diag_1            1666
diag_2            2894
diag_3            6481
number_diagnoses  0
max_glu_serum     0
A1Cresult         0
metformin         0
repaglinide       0
nateglinide       0
chlorpropamide    0
glimepiride       0
acetohexamide     0
glipizide         0
glyburide         0
tolbutamide       0
pioglitazone      0
rosiglitazone     0
acarbose         0
miglitol         0
troglitazone      0
tolazamide       0
examide          0
citoglipton       0
insulin          0
glyburide-metformin 0
glipizide-metformin 0
glimepiride-pioglitazone 0
metformin-rosiglitazone 0
metformin-pioglitazone 0
change           0
diabetesMed       0
readmitted       0
dtype: int64
encounter_id      int64
```



```

patient_nbr          int64
race                 object
gender               object
age                 object
weight              object
admission_type_id    int64
discharge_disposition_id int64
admission_source_id  int64
time_in_hospital     int64
payer_code           object
medical_specialty    object
num_lab_procedures   int64
num_procedures        int64
num_medications       int64
number_outpatient     int64
number_emergency      int64
number_inpatient      int64
diag_1               float64
diag_2               float64
diag_3               float64
number_diagnoses      int64
max_glu_serum        object
A1Cresult            object
metformin            object
repaglinide          object
nateglinide          object
chlorpropamide       object
glimepiride          object
acetohexamide        object
glipizide            object
glyburide            object
tolbutamide          object
pioglitazone         object
rosiglitazone        object
acarbose             object
miglitol             object
troglitazone         object
tolazamide           object
examide              object
citoglipton          object
insulin              object
glyburide-metformin  object
glipizide-metformin  object
glimepiride-pioglitazone object
metformin-rosiglitazone object
metformin-pioglitazone object
change               object
diabetesMed           object
readmitted           object
dtype: object

```

Out[25]:

	encounter_id	patient_nbr	race	gender	age	weight	admission_type_id	discharge_disposition_id	admission_source_id	time_in
0	2278392	8222157	Caucasian	Female	[0-10)	NaN	6	25	1	
1	149190	55629189	Caucasian	Female	[10-20)	NaN	1	1	7	
2	64410	86047875	AfricanAmerican	Female	[20-30)	NaN	1	1	7	
3	500364	82442376	Caucasian	Male	[30-40)	NaN	1	1	7	
4	16680	42519267	Caucasian	Male	[40-50)	NaN	1	1	7	

5 rows × 50 columns

time: 281 ms (started: 2022-12-25 21:11:30 +05:00)

In [26]:

```
diabetic_df = diabetic_df.sample(50000) #sampling rows from the dataset
```

time: 62 ms (started: 2022-12-25 21:11:31 +05:00)

In [27]:

```

#Dropping all the ID columns and columns with a lot of missing values
diabetic_classes = diabetic_df[['diabetesMed']]
diabetic_df.drop(columns = ['diabetesMed', 'encounter_id', 'patient_nbr', 'weight', 'admission_type_id', 'discharge_disposition_id', 'admission_source_id', 'payer_code', 'medical_specialty', 'encounter_id'], inplace = True) #Dropping unnecessary columns

#Replacing missing values by Nan
imputer = SimpleImputer(missing_values=np.nan, strategy='most_frequent')
imputer = imputer.fit(diabetic_df)

```

```
diabetic_df = pd.DataFrame(imputer.transform(diabetic_df), columns = (diabetic_df.columns)).astype(diabetic_df.dtypes)

#dummy-encoding (One-hot encoding) the categorical variables
diabetic_df = pd.get_dummies(diabetic_df, drop_first = True)
diabetic_df.shape

#Scaling and One hot Encoding
Scaler = MinMaxScaler()
diabetic_df = pd.get_dummies(diabetic_df)
diabetic_df = pd.DataFrame(Scaler.fit_transform(diabetic_df), columns = diabetic_df.columns)
print('Shape of df now is: ', diabetic_df.shape)
diabetic_df.head()
```

Shape of df now is: (50000, 77)

Out[27]:

	time_in_hospital	num_lab_procedures	num_procedures	num_medications	number_outpatient	number_emergency	number_inpatient	diag
0	0.15	0.00	0.00	0.16	0.00	0.00	0.00	0.
1	0.54	0.44	0.17	0.19	0.00	0.00	0.00	1.
2	0.38	0.11	0.50	0.23	0.11	0.00	0.00	0.
3	0.15	0.30	0.33	0.17	0.03	0.08	0.19	0.
4	0.77	0.40	0.50	0.33	0.00	0.00	0.00	0.

5 rows × 77 columns

time: 375 ms (started: 2022-12-25 21:11:34 +05:00)

In [28]:

```
diabetic_classes.loc[
    ((diabetic_classes['diabetesMed'] == 'Yes')), 'diabetesMed'] = 1    #Labeling the Yes case

diabetic_classes.loc[
    ((diabetic_classes['diabetesMed'] == 'No')), 'diabetesMed'] = 0    #Labeling the No case
```

time: 0 ns (started: 2022-12-25 21:11:37 +05:00)

5. Income Dataset

In [29]:

```
print(income_df.isnull().sum()) #Check missing values
print(income_df.dtypes)        #Check data types
income_df.head()              #Bird's eye view of dataset
```

```
row ID          0
Age             0
WorkClass       0
X1              0
Education Level 0
X2              0
Marital Status  0
Occupation      0
X3              0
Gender          0
X4              0
X5              0
Hours Per Week Working 0
Native Country  0
High Income     0
dtype: int64

row ID          object
Age             float64
WorkClass       int64
X1              float64
Education Level int64
X2              float64
Marital Status  int64
Occupation      int64
X3              int64
Gender          int64
X4              float64
X5              float64
Hours Per Week Working float64
Native Country  int64
High Income     int64
dtype: object
```

Out[29]:

	row ID	Age	WorkClass	X1	Education Level	X2	Marital Status	Occupation	X3	Gender	X4	X5	Hours Per Week Working	Native Country	High Income
0	Row2	38.00	2	215646.00	1	9.00	2	2	0	0	0.00	0.00	40.00	0	0
1	Row3	53.00	2	234721.00	2	7.00	1	2	1	0	0.00	0.00	40.00	0	0
2	Row5	37.00	2	284582.00	3	14.00	1	1	0	1	0.00	0.00	40.00	0	0
3	Row7	52.00	1	209642.00	1	9.00	1	1	0	0	0.00	0.00	45.00	0	1
4	Row8	31.00	2	45781.00	3	14.00	0	3	0	1	14084.00	0.00	50.00	0	1

time: 16 ms (started: 2022-12-25 21:11:39 +05:00)

In [30]:

```
income_df = income_df.sample(40000) #sampling rows from the dataset
```

time: 16 ms (started: 2022-12-25 21:11:39 +05:00)

In [31]:

```
income_classes = income_df[['High Income']]
income_df.drop(columns = ['High Income', 'row ID'], inplace = True) #Dropping unnecessary columns

#Scaling and One hot Encoding
Scaler = MinMaxScaler()
income_df = pd.get_dummies(income_df)
income_df = pd.DataFrame(Scaler.fit_transform(income_df), columns = income_df.columns)
print('Shape of df now is: ', income_df.shape)
income_df.head()
```

Shape of df now is: (40000, 13)

Out[31]:

	Age	WorkClass	X1	Education Level	X2	Marital Status	Occupation	X3	Gender	X4	X5	Hours Per Week Working	Native Country
0	0.15	0.62	0.09	0.07	0.53	0.00	0.79	0.00	1.00	0.01	0.00	0.32	0.00
1	0.58	0.25	0.24	0.27	0.27	0.17	0.29	0.00	0.00	0.00	0.00	0.30	0.00
2	0.29	0.75	0.47	0.67	0.53	0.67	0.21	0.25	0.00	0.24	0.14	0.24	0.00
3	0.10	0.25	0.21	0.80	0.47	0.00	0.29	0.25	0.00	0.00	0.00	0.32	0.00
4	0.25	0.25	0.06	0.07	0.74	0.00	0.43	0.00	0.00	0.02	0.00	0.42	0.00

time: 16 ms (started: 2022-12-25 21:11:40 +05:00)

6. Dry Beans Dataset

In [308]:

```
print(beans_df.isnull().sum()) #Check missing values
print(beans_df.dtypes) #Check data types
beans_df.head() #Bird's eye view of dataset
```

```
Area          0
Perimeter     0
MajorAxisLength  0
MinorAxisLength  0
AspectRatio    0
Eccentricity   0
ConvexArea     0
EquivDiameter  0
Extent         0
Solidity       0
roundness      0
Compactness    0
ShapeFactor1   0
ShapeFactor2   0
ShapeFactor3   0
ShapeFactor4   0
Class          0
dtype: int64
Area          int64
Perimeter     float64
MajorAxisLength  float64
MinorAxisLength  float64
AspectRatio    float64
Eccentricity   float64
ConvexArea     int64
EquivDiameter  float64
Extent         float64
Solidity       float64
roundness      float64
Compactness    float64
ShapeFactor1   float64
ShapeFactor2   float64
```

```
ShapeFactor3      float64
ShapeFactor4      float64
Class             object
dtype: object
```

```
Out[308...
   Area  Perimeter  MajorAxisLength  MinorAxisLength  AspectRatio  Eccentricity  ConvexArea  EquivDiameter  Extent  Solidity  roundness
0  28395    610.29         208.18         173.89         1.20         0.55        28715         190.14      0.76     0.99     0.96
1  28734    638.02         200.52         182.73         1.10         0.41        29172         191.27      0.78     0.98     0.85
2  29380    624.11         212.83         175.93         1.21         0.56        29690         193.41      0.78     0.99     0.95
3  30008    645.88         210.56         182.52         1.15         0.50        30724         195.47      0.78     0.98     0.90
4  30140    620.13         201.85         190.28         1.06         0.33        30417         195.90      0.77     0.99     0.98
```

time: 16 ms (started: 2022-12-29 01:01:20 +05:00)

```
In [309...
beans_classes = beans_df[['Class']]
beans_df.drop(columns = ['Class'], inplace = True) #Dropping unnecessary columns

#Label encoding the Y variable
le = LabelEncoder()
beans_classes = pd.DataFrame(le.fit_transform(beans_classes), columns = beans_classes.columns)

#Scaling and One hot Encoding
Scaler = MinMaxScaler()
beans_df = pd.get_dummies(beans_df)
beans_df = pd.DataFrame(Scaler.fit_transform(beans_df), columns = beans_df.columns)
print('Shape of df now is: ', beans_df.shape)
beans_df.head()
```

Shape of df now is: (13611, 16)

```
Out[309...
   Area  Perimeter  MajorAxisLength  MinorAxisLength  AspectRatio  Eccentricity  ConvexArea  EquivDiameter  Extent  Solidity  roundness
0   0.03      0.06         0.04         0.15         0.12         0.48         0.03         0.07      0.67     0.92     0.93
1   0.04      0.08         0.03         0.18         0.05         0.28         0.03         0.07      0.74     0.87     0.79
2   0.04      0.07         0.05         0.16         0.13         0.50         0.04         0.08      0.72     0.93     0.91
3   0.04      0.08         0.05         0.18         0.09         0.40         0.04         0.08      0.73     0.76     0.83
4   0.04      0.07         0.03         0.20         0.03         0.17         0.04         0.08      0.70     0.95     0.99
```

time: 16 ms (started: 2022-12-29 01:01:26 +05:00)

7. Bank Notes Detection Dataset

```
In [85]:
print(banknotes_df.isnull().sum()) #Check missing values
print(banknotes_df.dtypes) #Check data types
banknotes_df.head() #Bird's eye view of dataset
```

```
variance      0
skewness      0
curtosis      0
entropy       0
Class         0
dtype: int64
variance      float64
skewness      float64
curtosis      float64
entropy       float64
Class         int64
dtype: object
```

```
Out[85]:
   variance  skewness  curtosis  entropy  Class
0     3.62     8.67    -2.81    -0.45     0
1     4.55     8.17    -2.46    -1.46     0
2     3.87    -2.64     1.92     0.11     0
3     3.46     9.52    -4.01    -3.59     0
4     0.33    -4.46     4.57    -0.99     0
```

time: 0 ns (started: 2022-12-26 01:52:46 +05:00)

```
In [86]:
banknotes_classes = banknotes_df[['Class']]
banknotes_df.drop(columns = ['Class'], inplace = True) #Dropping unnecessary columns
```

```
#Scaling and One hot Encoding
Scaler = MinMaxScaler()
banknotes_df = pd.get_dummies(banknotes_df)
banknotes_df = pd.DataFrame(Scaler.fit_transform(banknotes_df), columns = banknotes_df.columns)
print('Shape of df now is: ', banknotes_df.shape)
banknotes_df.head()
```

Shape of df now is: (1372, 4)

Out[86]:

	variance	skewness	curtosis	entropy
0	0.77	0.84	0.11	0.74
1	0.84	0.82	0.12	0.64
2	0.79	0.42	0.31	0.79
3	0.76	0.87	0.05	0.45
4	0.53	0.35	0.42	0.69

time: 15 ms (started: 2022-12-26 01:52:47 +05:00)

8. Audit Risk Dataset

In [87]:

```
print(audit_df.isnull().sum()) #Check missing values
print(audit_df.dtypes) #Check data types
audit_df.head() #Bird's eye view of dataset
```

```
Sector_score      0
LOCATION_ID        0
PARA_A           0
Score_A          0
Risk_A           0
PARA_B           0
Score_B          0
Risk_B           0
TOTAL            0
numbers          0
Score_B.1        0
Risk_C           0
Money_Value      1
Score_MV         0
Risk_D           0
District_Loss    0
PROB             0
Risk_E           0
History          0
Prob            0
Risk_F           0
Score            0
Inherent_Risk    0
CONTROL_RISK     0
Detection_Risk   0
Audit_Risk       0
Risk             0
dtype: int64
Sector_score      float64
LOCATION_ID        object
PARA_A           float64
Score_A          float64
Risk_A           float64
PARA_B           float64
Score_B          float64
Risk_B           float64
TOTAL            float64
numbers          float64
Score_B.1        float64
Risk_C           float64
Money_Value      float64
Score_MV         float64
Risk_D           float64
District_Loss    int64
PROB             float64
Risk_E           float64
History          int64
Prob            float64
Risk_F           float64
Score            float64
Inherent_Risk    float64
CONTROL_RISK     float64
Detection_Risk   float64
Audit_Risk       float64
Risk             int64
dtype: object
```

Out[87]:

	Sector_score	LOCATION_ID	PARA_A	Score_A	Risk_A	PARA_B	Score_B	Risk_B	TOTAL	numbers	...	RiSk_E	History	Prob	Risk_F	Sci
0	3.89	23	4.18	0.60	2.51	2.50	0.20	0.50	6.68	5.00	...	0.40	0	0.20	0.00	2
1	3.89	6	0.00	0.20	0.00	4.83	0.20	0.97	4.83	5.00	...	0.40	0	0.20	0.00	2
2	3.89	6	0.51	0.20	0.10	0.23	0.20	0.05	0.74	5.00	...	0.40	0	0.20	0.00	2
3	3.89	6	0.00	0.20	0.00	10.80	0.60	6.48	10.80	6.00	...	0.40	0	0.20	0.00	4
4	3.89	6	0.00	0.20	0.00	0.08	0.20	0.02	0.08	5.00	...	0.40	0	0.20	0.00	2

5 rows × 27 columns

time: 16 ms (started: 2022-12-26 01:52:49 +05:00)

In [88]:

```
audit_df = audit_df.dropna()
audit_classes = audit_df[['Risk']]
audit_df.drop(columns = ['Risk', 'LOCATION_ID'], inplace = True) #Dropping unnecessary columns

#Scaling and One hot Encoding
Scaler = MinMaxScaler()
audit_df = pd.get_dummies(audit_df)
audit_df = pd.DataFrame(Scaler.fit_transform(audit_df), columns = audit_df.columns)
print('Shape of df now is: ', audit_df.shape)
audit_df.head()
```

Shape of df now is: (775, 25)

Out[88]:

	Sector_score	PARA_A	Score_A	Risk_A	PARA_B	Score_B	Risk_B	TOTAL	numbers	Score_B.1	...	PROB	RiSk_E	History	Prob	Risk_F
0	0.04	0.05	1.00	0.05	0.00	0.00	0.00	0.01	0.00	0.00	...	0.00	0.00	0.00	0.00	0.00
1	0.04	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	...	0.00	0.00	0.00	0.00	0.00
2	0.04	0.01	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	...	0.00	0.00	0.00	0.00	0.00
3	0.04	0.00	0.00	0.00	0.01	1.00	0.01	0.01	0.25	1.00	...	0.00	0.00	0.00	0.00	0.00
4	0.04	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	...	0.00	0.00	0.00	0.00	0.00

5 rows × 25 columns

time: 15 ms (started: 2022-12-26 01:57:02 +05:00)

5.2.3 Dimensionality Reduction Techniques with ML

A) Dimensionality Reduction Pipeline for Classification datasets:

In [313...]

```
def dimensionality_reduction_classification(X_train, y_train, X_test, y_test, pca_dim = 0.95, ipca_dim= 0.95, svd_dim

    print("Starting DR Pipeline...")

    print("1. Running Lazy Predict without DR")

    #Running Lazy Predict without Dimensionality Reduction:
    clf = LazyClassifier(verbose=0, ignore_warnings=True, custom_metric = None)
    simple_models = clf.fit(X_train, X_test, y_train, y_test)[0].sort_index()
    simple_models['dim'] = X_train.shape[1]
    simple_models.drop(columns = ['Balanced Accuracy', 'Time Taken'], inplace = True)

    print("Success!")

    #Model Performances with PCA:

    #Running PCA:
    print("2. Running PCA")
    pca = PCA(n_components = pca_dim)
    pca.fit(X_train)
    X_train_transformed = pca.transform(X_train)
    X_test_transformed = pca.transform(X_test)
    print("Success!")

    #Running Lazy Predict with PCA
    print("3. Running Lazy Predict on PCA dataset")
    clf = LazyClassifier(verbose=0, ignore_warnings=True, custom_metric = None)
    pcamodels = clf.fit(X_train_transformed, X_test_transformed, y_train, y_test)[0].sort_index()
    pcamodels.drop(columns = ['Balanced Accuracy', 'Time Taken'], inplace = True)
    pcamodels['dims'] = len(pca.components_)
    print("Success!")

    #Model Performances with Incremental-PCA:
```

```

#Running Incremental PCA:
print("4. Running Incremental PCA")
for i in range(1, X_train.shape[1], 1):
    ipca = IncrementalPCA(n_components = i)
    ipca.fit(X_train)
    X_train_transformed = ipca.transform(X_train)
    X_test_transformed = ipca.transform(X_test)

    if ipca.explained_variance_ratio_.sum() >= ipca_dim:
        print("Success!")
        print("5. Running Lazy Predict on Incremental PCA dataset")
        clf = LazyClassifier(verbose=0, ignore_warnings=True, custom_metric = None) #Running Lazy Predict after
        ipcamodels = clf.fit(X_train_transformed, X_test_transformed, y_train, y_test)[0].sort_index()
        ipcamodels.drop(columns = ['Balanced Accuracy', 'Time Taken'], inplace = True)
        ipcamodels['dims'] = ipca.n_components_
        print("Success!")
        break

#Model Performances with Sparse-PCA:
print("6. Running Sparse PCA")
spca = SparsePCA(n_components = 10)
spca.fit(X_train)
X_train_transformed = spca.transform(X_train)
X_test_transformed = spca.transform(X_test)
print("Success!")

#Running Lazy Predict Sparse PCA:
print("7. Running Lazy Predict on Sparse PCA dataset")
clf = LazyClassifier(verbose=0, ignore_warnings=True, custom_metric = None) #Running Lazy Predict after SVD
spcamodels = clf.fit(X_train_transformed, X_test_transformed, y_train, y_test)[0].sort_index()
spcamodels.drop(columns = ['Balanced Accuracy', 'Time Taken'], inplace = True)
spcamodels['dims'] = spca.n_components_
print("Success!")

#Model Performances with LDA:

#Running LDA:
print("8. Running LDA")
lda = LinearDiscriminantAnalysis()
lda.fit(X_train, y_train)
X_train_transformed = lda.transform(X_train)
X_test_transformed = lda.transform(X_test)
print("Success!")

#Running Lazy Predict with LDA
print("9. Running Lazy Predict on LDA dataset")
clf = LazyClassifier(verbose=0, ignore_warnings=True, custom_metric = None)
ldamodels = clf.fit(X_train_transformed, X_test_transformed, y_train, y_test)[0].sort_index()
ldamodels.drop(columns = ['Balanced Accuracy', 'Time Taken'], inplace = True)
ldamodels['dims'] = len(lda.coef_)
print("Success!")

#Model Performances with SVD:

#Running SVD:
print("10. Running SVD")
for i in range(1, X_train.shape[1], 1):
    svd = TruncatedSVD(n_components = i)
    svd.fit(X_train)
    X_train_transformed = svd.transform(X_train)
    X_test_transformed = svd.transform(X_test)

    if svd.explained_variance_ratio_.sum() >= svd_dim or i>=X_train.shape[1]-1:
        print("Success!")
        print("11. Running Lazy Predict on SVD dataset")
        clf = LazyClassifier(verbose=0, ignore_warnings=True, custom_metric = None) #Running Lazy Predict after
        svdmodels = clf.fit(X_train_transformed, X_test_transformed, y_train, y_test)[0].sort_index()
        svdmodels.drop(columns = ['Balanced Accuracy', 'Time Taken'], inplace = True)
        svdmodels['dims'] = len(svd.components_)
        print("Success!")
        break

#Compiling Model Results:
print("Compiling Model Results")
models_results = pd.concat([simple_models,
                             pcamodels,
                             ipcamodels,
                             spcamodels,
                             ldamodels,
                             svdmodels], axis = 1, keys = ['Without DR', 'PCA ', 'Incremental-PCA', 'Sparse-PCA', 'LDA', 'SVD'])

print("Pipeline run Successful")

```

Out[13]:																
Without DR								PCA	Incremental-PCA	...	Sparse-PCA					
	Accuracy	ROC AUC	F1 Score	dim	Accuracy	ROC AUC	F1 Score	dims	Accuracy	ROC AUC	...	F1 Score	dims	Accuracy	ROC AUC	F1 Score
Model																
AdaBoostClassifier	0.87	0.68	0.85	35	0.85	0.63	0.83	18	0.87	0.66	...	0.84	10	0.86	0.68	0.85
BernoulliNB	0.78	0.68	0.79	35	0.84	0.56	0.79	18	0.84	0.58	...	0.84	10	0.83	0.50	0.75
DecisionTreeClassifier	0.84	0.70	0.84	35	0.84	0.67	0.83	18	0.82	0.65	...	0.80	10	0.83	0.66	0.82

Model	Without DR							PCA	Incremental-PCA		...	Sparse-PCA					
	Accuracy	ROC AUC	F1 Score	dim	Accuracy	ROC AUC	F1 Score	dims	Accuracy	ROC AUC	...	F1 Score	dims	Accuracy	ROC AUC	F1 Score	
KNeighborsClassifier	0.85	0.61	0.82	35	0.84	0.61	0.82	18	0.84	0.61	...	0.83	10	0.85	0.66	0.83	
LinearSVC	0.86	0.65	0.84	35	0.87	0.65	0.85	18	0.88	0.65	...	0.84	10	0.87	0.64	0.84	
LogisticRegression	0.88	0.68	0.86	35	0.87	0.66	0.85	18	0.87	0.66	...	0.86	10	0.86	0.64	0.84	
RandomForestClassifier	0.86	0.62	0.83	35	0.86	0.64	0.83	18	0.86	0.63	...	0.84	10	0.83	0.66	0.82	
XGBClassifier	0.88	0.71	0.87	35	0.87	0.69	0.85	18	0.86	0.66	...	0.84	10	0.86	0.65	0.84	

```
time: 31 ms (started: 2022-12-31 02:39:22 +05:00)
```

1. For the marketing dataset, initially with full dataset having 35 features gave us the maximum accuracy of 0.88 and F1 score of 0.87 from the XG Boost Algorithm. For Naive bayes, we saw an accuracy of 0.78 and F1 score of 0.79.
2. After applying normal PCA, the accuracy and F1 score are almost the same as without applying PCA which means that almost all variation was captured by the Principal Components. Furthermore, Naive Bayes algorithm performance increased greatly when feature reduction was applied as its performance increased from 0.76 to 0.84. However, it is important to note that AUC-ROC of the algorithm dropped.
3. Applying the other PCA variants such as Incremental and Sparse PCA does not improve the results and they perform at Par with PCA although sparse PCA reduced the dimensions further to 10 with truly little sacrifice to variance capture. Naive Bayes algorithm's performance increased further as well. However, it is important to note that AUC-ROC of the algorithm dropped.
4. Then Linear Discriminant Analysis was tried, and the performance was at Par with the other DR techniques. However, since it reduces the dimensions to 1, it proves to be the best DR technique so far.
5. Lastly, Singular Value Decomposition also performed at PAR with the other DR techniques and gave promising results.

2. Credit Card Dataset

```
time: 16 ms (started: 2022-12-25 18:35:54 +05:00)
```

Starting DR Pipeline...

[illegible]

2. Running PCA

3. Running Lazy Predict on PCA dataset

Success!

4. Running Incremental PCA

5. Running Lazy Predict on Incremental PCA dataset

Success!

6. Running Sparse PCA


```
Starting DR Pipeline...
1. Running Lazy Predict without DR
100%|██████████████████████████████████████████████████████████████████████████| 29/29 [04:29<00:00, 9.29s/i]
t]
Success!
2. Running PCA
Success!
3. Running Lazy Predict on PCA dataset
100%|██████████████████████████████████████████████████████████████████████████| 29/29 [04:24<00:00, 9.12s/i]
t]
Success!
4. Running Incremental PCA
Success!
5. Running Lazy Predict on Incremental PCA dataset
100%|██████████████████████████████████████████████████████████████████████████| 29/29 [04:18<00:00, 8.92s/i]
t]
Success!
6. Running Sparse PCA
Success!
7. Running Lazy Predict on Sparse PCA dataset
100%|██████████████████████████████████████████████████████████████████████████| 29/29 [03:32<00:00, 7.34s/i]
t]
Success!
8. Running LDA
Success!
9. Running Lazy Predict on LDA dataset
100%|██████████████████████████████████████████████████████████████████████████| 29/29 [02:58<00:00, 6.16s/i]
t]
Success!
10. Running SVD
Success!
11. Running Lazy Predict on SVD dataset
100%|██████████████████████████████████████████████████████████████████████████| 29/29 [04:25<00:00, 9.14s/i]
t]
Success!
Compiling Model Results
Pipeline run Successful
time: 24min 22s (started: 2022-12-25 18:55:49 +05:00)
```

Since the heart disease dataset is slightly imbalanced, we can use AUC-ROC or F1 score to gauge its performance. I would be making comparison based on both of these.

1. Without DR, the best F1-score is achieved to be 0.89 by several algorithms with an AUC-ROC of 0.56. This would be used as a benchmark for further comparisons.
2. When PCA is applied, the F1-score and AUC-ROC were still the same as without DR which proved to be a good DR technique as it reduces the dims from 50 to 28.
3. Similarly, for other PCA variants, they all performed at PAR however, sparse PCA performs better in a way that it does not compromise on variance capture and reduces dimensions the most which is 10.
4. When LDA is applied, there is barely any compromise on variance capture as the metrics remain same however, the dimensions decrease to 1 which is its major achievement.
5. SVD performs at par with the other techniques however, the dims are still too high compared to the others.

It is becoming evident that LDA is the winner DR technique!

4. Diabetes Dataset

```
In [32]: #Lets split the data into a train test split from the start, test set will be kept separate and will only be used for  
X_train, X_test, y_train, y_test = train_test_split(diabetic_df, diabetic_classes.astype('int'), test_size=0.25, rand  
time: 16 ms (started: 2022-12-25 21:11:50 +05:00)
```

```
In [33]: #Pipeline run
models_results = dimensionality_reduction_classification(X_train, y_train, X_test, y_test, pca_dim = 0.95, svd_dim = 0.95)

#keeping only the desired algorithms
results = models_results.T[['AdaBoostClassifier', 'BernoulliNB',
                             'DecisionTreeClassifier', 'KNeighborsClassifier', 'LinearSVC', 'LogisticRegression',
                             'RandomForestClassifier', 'XGBClassifier']].T

#Exporting results to excel
results.to_excel('Results/Classification/Diabetes.xlsx', sheet_name = 'Diabetes Dataset')
```

```
Starting DR Pipeline...
1. Running Lazy Predict without DR
100%|██████████████████████████████████████████████████████████████████████████████| 29/29 [03:43<00:00, 7.72s/i]
Success!
2. Running PCA
Success!
3. Running Lazy Predict on PCA dataset
100%|██████████████████████████████████████████████████████████████████████████████| 29/29 [03:19<00:00, 6.87s/i]
Success!
4. Running Incremental PCA
Success!
5. Running Lazy Predict on Incremental PCA dataset
100%|██████████████████████████████████████████████████████████████████████████████| 29/29 [03:19<00:00, 6.89s/i]
Success!
6. Running Sparse PCA
Success!
7. Running Lazy Predict on Sparse PCA dataset
100%|██████████████████████████████████████████████████████████████████████████████| 29/29 [03:00<00:00, 6.23s/i]
Success!
8. Running LDA
Success!
9. Running Lazy Predict on LDA dataset
100%|██████████████████████████████████████████████████████████████████████████████| 29/29 [03:03<00:00, 6.33s/i]
Success!
10. Running SVD
Success!
11. Running Lazy Predict on SVD dataset
100%|██████████████████████████████████████████████████████████████████████████████| 29/29 [03:23<00:00, 7.03s/i]
Success!
Compiling Model Results
Pipeline run Successful
time: 20min 51s (started: 2022-12-25 21:11:50 +05:00)
```

Dataset Results:

```
In [301... #Results
results = pd.read_excel('Results/Classification/diabetes.xlsx', header=[0, 1], index_col=0 )
results
```

	Without DR			PCA			Incremental-PCA			Sparse-PC						
	Accuracy	ROC AUC	F1 Score	dim	Accuracy	ROC AUC	F1 Score	dims	Accuracy	ROC AUC	F1 Score	dims	Accuracy	ROC AUC	F1 Score	dim
Model																
AdaBoostClassifier	1.00	1.00	1.00	77	1.00	1.00	1.00	30	1.00	1.00	1.00	30	0.99	0.99	0.99	1
BernoulliNB	1.00	1.00	1.00	77	0.96	0.95	0.96	30	0.94	0.90	0.93	30	0.93	0.95	0.94	1
DecisionTreeClassifier	1.00	1.00	1.00	77	1.00	1.00	1.00	30	1.00	0.99	1.00	30	0.98	0.98	0.98	1
KNeighborsClassifier	0.99	0.99	0.99	77	0.99	0.99	0.99	30	0.99	0.99	0.99	30	0.99	0.99	0.99	1
LinearSVC	1.00	1.00	1.00	77	1.00	1.00	1.00	30	1.00	1.00	1.00	30	0.99	0.99	0.99	1
LogisticRegression	1.00	1.00	1.00	77	1.00	1.00	1.00	30	1.00	1.00	1.00	30	0.99	0.99	0.99	1
RandomForestClassifier	1.00	1.00	1.00	77	1.00	1.00	1.00	30	1.00	1.00	1.00	30	0.99	0.99	0.99	1
XGBClassifier	1.00	1.00	1.00	77	1.00	1.00	1.00	30	1.00	1.00	1.00	30	0.99	0.99	0.99	1

Analysis:

5. Income Dataset

```
In [38]: #Pipeline run
models_results = dimensionality_reduction_classification(X_train, y_train, X_test, y_test, pca_dim = 0.95, svd_dim = 0.95)

#keeping only the desired algorithms
results = models_results.T[['AdaBoostClassifier', 'BernoulliNB',
                             'DecisionTreeClassifier', 'KNeighborsClassifier', 'LinearSVC', 'LogisticRegression',
                             'RandomForestClassifier', 'XGBClassifier']].T

#Exporting results to excel
results.to_excel('Results/Classification/Income.xlsx', sheet_name = 'Income Dataset')
```

```
Starting DR Pipeline...
1. Running Lazy Predict without DR
100%|██████████████████████████████████████████████████████████████████████████| 29/29 [02:53<00:00, 5.99s/i
t]
Success!
2. Running PCA
Success!
3. Running Lazy Predict on PCA dataset
100%|██████████████████████████████████████████████████████████████████████████| 29/29 [02:58<00:00, 6.15s/i
t]
Success!
4. Running Incremental PCA
Success!
5. Running Lazy Predict on Incremental PCA dataset
100%|██████████████████████████████████████████████████████████████████████████| 29/29 [03:01<00:00, 6.25s/i
t]
Success!
6. Running Sparse PCA
Success!
7. Running Lazy Predict on Sparse PCA dataset
100%|██████████████████████████████████████████████████████████████████████████| 29/29 [02:48<00:00, 5.80s/i
t]
Success!
8. Running LDA
Success!
9. Running Lazy Predict on LDA dataset
100%|██████████████████████████████████████████████████████████████████████████| 29/29 [02:17<00:00, 4.75s/i
t]
Success!
10. Running SVD
Success!
11. Running Lazy Predict on SVD dataset
```

```
100%|██████████████████████████████████████████████████████████████████████████| 29/29 [03:00<00:00, 6.23s/i  
t]  
Success!  
Compiling Model Results  
Pipeline run Successful  
time: 17min 4s (started: 2022-12-25 21:55:32 +05:00)
```

Dataset Results:

In [302...

```
#Results
results = pd.read_excel('Results/Classification/income.xlsx', header=[0, 1], index_col=0)
results
```

Out[302...]

Model	Without DR				PCA				Incremental-PCA				Sparse-PC			
	Accuracy	ROC AUC	F1 Score	dim	Accuracy	ROC AUC	F1 Score	dims	Accuracy	ROC AUC	F1 Score	dims	Accuracy	ROC AUC	F1 Score	dim
AdaBoostClassifier	0.75	0.69	0.74	13	0.72	0.64	0.71	12	0.71	0.63	0.70	12	0.74	0.67	0.73	1
BernoulliNB	0.69	0.63	0.68	13	0.70	0.57	0.65	12	0.69	0.56	0.64	12	0.69	0.62	0.68	1
DecisionTreeClassifier	0.75	0.72	0.75	13	0.72	0.68	0.72	12	0.72	0.68	0.72	12	0.73	0.69	0.73	1
KNeighborsClassifier	0.74	0.69	0.74	13	0.74	0.69	0.74	12	0.74	0.69	0.74	12	0.74	0.69	0.73	1
LinearSVC	0.69	0.57	0.65	13	0.69	0.57	0.65	12	0.69	0.57	0.65	12	0.69	0.57	0.65	1
LogisticRegression	0.69	0.57	0.65	13	0.69	0.57	0.65	12	0.69	0.57	0.65	12	0.69	0.57	0.65	1
RandomForestClassifier	0.78	0.75	0.78	13	0.76	0.70	0.75	12	0.77	0.71	0.76	12	0.76	0.72	0.76	1
XGBClassifier	0.77	0.73	0.77	13	0.75	0.70	0.74	12	0.76	0.71	0.75	12	0.75	0.71	0.75	1

time: 15 ms (started: 2022-12-29 00:59:52 +05:00)

Analysis:

The dataset is not that imbalanced so we can use any of the three metrics to compare the DR techniques. To keep it coherent, lets focus on F1 score for this dataset:

1. The best F1 score achieved with full dataset without any DR technique is around 0.78 given by Random Forest.
2. After applying PCA, a very slight drop in F1 score is seen however, the dims are not reduced that much either and remain at 12 reduced from 13. So, at the compromise of little variance capture, only 1 feature is reduced.
3. Applying different variants of PCA resulted in comparable results. However, sparse PCA performs better in a way that it reduces dimensions from 12 to 10 with truly little compromise on variance capture.
4. After applying LDA, the features are reduced to 1 however, there is a significant drop in variance capture which shows that LDA does not perform good on this dataset.
5. Lastly, SVD performed like PCA.

It is seen that in this dataset, DR techniques cause a significant decrease in variance capture if features are reduced to less and the conclusion reached from this is that the dataset already contains the most important information.

6. Dry Beans Dataset

In [314...

```
# Lets split the data into a train test split from the start, test set will be kept separate and will only be used for
X_train, X_test, y_train, y_test = train_test_split(beans_df, beans_classes.astype('int'), test_size=0.25, random_state=42)
time: 0 ns (started: 2022-12-29 01:02:12 +05:00)
```

In [315...

```
#Pipeline run
models_results = dimensionality_reduction_classification(X_train, y_train, X_test, y_test, pca_dim = 0.95, svd_dim = 0.95)

#keeping only the desired algorithms
results = results[[]]
results = models_results.T[['AdaBoostClassifier', 'BernoulliNB',
                            'DecisionTreeClassifier', 'KNeighborsClassifier', 'LinearSVC', 'LogisticRegression',
                            'RandomForestClassifier', 'XGBClassifier']].T

#Exporting results to excel
results.to_excel('Results/Classification/Beans.xlsx', sheet name = 'Beans Dataset')
```

Starting DR Pipeline...

1. Running Lazy Predict without DR

```
100%|██████████████████████████████████████████████████████████████████████████| 29/29 [00:19<00:00, 1.48it/
s]
Success!
2. Running PCA
```

Dataset Results:

```
#Results
results = pd.read_excel('Results/Classification/Beans.xlsx', header=[0, 1], index_col=0)
results
```

	Without DR			PCA			Incremental-PCA			Sparse-PCA			LD	
	Accuracy	F1 Score	dim	Accuracy	F1 Score	dims	Accuracy	F1 Score	dims	Accuracy	F1 Score	dims		
Model														
AdaBoostClassifier	0.68	0.62	16	0.44	0.27	4	0.44	0.27	4	0.74	0.70	10	0.73	0.72
BernoulliNB	0.72	0.72	16	0.65	0.58	4	0.66	0.58	4	0.77	0.76	10	0.85	0.85
DecisionTreeClassifier	0.89	0.89	16	0.85	0.85	4	0.85	0.85	4	0.89	0.89	10	0.89	0.89
KNeighborsClassifier	0.93	0.93	16	0.89	0.89	4	0.89	0.89	4	0.93	0.93	10	0.92	0.92
LinearSVC	0.91	0.91	16	0.85	0.85	4	0.86	0.85	4	0.91	0.91	10	0.91	0.91
LogisticRegression	0.92	0.92	16	0.89	0.89	4	0.89	0.89	4	0.92	0.92	10	0.92	0.92
RandomForestClassifier	0.92	0.92	16	0.89	0.89	4	0.89	0.89	4	0.93	0.93	10	0.93	0.93
XGBClassifier	0.92	0.92	16	0.89	0.89	4	0.89	0.89	4	0.92	0.92	10	0.92	0.92

```
time: 15 ms (started: 2022-12-29 01:16:14 +05:00)
```

For the dry bean's dataset, AUC-ROC was removed since it is a multi-class problem. Since the dataset is balanced, I will be focusing on Accuracy as the metric for comparison.

- LDA still proves to be good as it maintains the best accuracy with significant reduction of features. The other DR techniques are still good enough**

7. Bank Notes Dataset

				Without DR				PCA		Incremental-PCA				Sparse-PC			
	Accuracy	ROC AUC	F1 Score	dim	Accuracy	ROC AUC	F1 Score	dims	Accuracy	ROC AUC	F1 Score	dims	Accuracy	ROC AUC	F1 Score	dim	
Model																	
AdaBoostClassifier	0.99	0.99	0.99	4	0.92	0.91	0.92	3	0.93	0.93	0.93	3	0.99	0.99	0.99	1	
BernoulliNB	0.86	0.86	0.86	4	0.79	0.77	0.78	3	0.79	0.77	0.78	3	0.86	0.86	0.86	1	
DecisionTreeClassifier	0.98	0.98	0.98	4	0.95	0.95	0.95	3	0.94	0.94	0.94	3	0.99	0.99	0.99	1	
KNeighborsClassifier	1.00	1.00	1.00	4	0.97	0.97	0.97	3	0.97	0.97	0.97	3	1.00	1.00	1.00	1	
LinearSVC	0.99	0.99	0.99	4	0.92	0.91	0.92	3	0.92	0.92	0.92	3	0.99	0.99	0.99	1	
LogisticRegression	0.98	0.98	0.98	4	0.92	0.92	0.92	3	0.92	0.92	0.92	3	0.98	0.98	0.98	1	
RandomForestClassifier	0.99	0.99	0.99	4	0.95	0.95	0.95	3	0.95	0.95	0.95	3	0.99	0.99	0.99	1	

Model	Without DR				PCA			Incremental-PCA				Sparse-PC.				
	Accuracy	ROC AUC	F1 Score	dim	Accuracy	ROC AUC	F1 Score	dims	Accuracy	ROC AUC	F1 Score	dims	Accuracy	ROC AUC	F1 Score	dim
XGBClassifier	0.99	0.99	0.99	4	0.94	0.95	0.94	3	0.94	0.95	0.94	3	0.99	0.99	0.99	1

Analysis:

1. The accuracy achieved without DR techniques is 1 which seems that the model is too good to be true.
2. After applying PCA, the accuracy is slightly compromised to 0.97 with just 1 feature reduction.
3. Different PCA variants perform similar with sparse PCA performing worse as it increases the dimensions from 4 to 10.
4. LDA captures all the variance of the dataset and does not compromise on accuracy while reducing the number of dimensions from 4 to 1 which shows that it performs very well.
5. SVD performs at par with PCA.

8. Audit Risk Dataset

```
In [90]: #Pipeline run
models_results = dimensionality_reduction_classification(X_train, y_train, X_test, y_test, pca_dim = 0.95, svd_dim = 0.95)

#keeping only the desired algorithms
results = models_results.T[['AdaBoostClassifier', 'BernoulliNB',
                             'DecisionTreeClassifier', 'KNeighborsClassifier', 'LinearSVC', 'LogisticRegression',
                             'RandomForestClassifier', 'XGBClassifier']].T

#Exporting results to excel
results.to_excel('Results/Classification/AuditRisk.xlsx', sheet_name = 'Audit Risk Dataset')
```

[illegible]

Success!
Compiling Model Results
Pipeline run Successful
time: 3.86 s (started: 2022-12-26 01:57:14 +05:00)

Dataset Results:

In [335...

```
#Results
results = pd.read_excel('Results/Classification/AuditRisk.xlsx', header=[0, 1], index_col=0)
results
```

Out[335...

Model	Without DR				PCA				Incremental-PCA				Sparse-PCA			
	Accuracy	ROC AUC	F1 Score	dim	Accuracy	ROC AUC	F1 Score	dims	Accuracy	ROC AUC	F1 Score	dims	Accuracy	ROC AUC	F1 Score	dim
AdaBoostClassifier	1.00	1.00	1.00	25	0.98	0.98	0.98	7	0.98	0.98	0.98	7	0.99	0.99	0.99	1
BernoulliNB	0.92	0.90	0.92	25	0.87	0.85	0.86	7	0.86	0.84	0.85	7	0.89	0.88	0.89	1
DecisionTreeClassifier	1.00	1.00	1.00	25	0.98	0.98	0.98	7	0.99	0.99	0.99	7	0.98	0.98	0.98	1
KNeighborsClassifier	0.97	0.96	0.97	25	0.96	0.95	0.96	7	0.96	0.95	0.96	7	0.96	0.95	0.96	1
LinearSVC	0.98	0.98	0.98	25	0.96	0.95	0.96	7	0.95	0.95	0.95	7	0.98	0.98	0.98	1
LogisticRegression	0.98	0.98	0.98	25	0.95	0.95	0.95	7	0.95	0.95	0.95	7	0.98	0.98	0.98	1
RandomForestClassifier	1.00	1.00	1.00	25	0.98	0.98	0.98	7	0.98	0.98	0.98	7	0.98	0.98	0.98	1
XGBClassifier	1.00	1.00	1.00	25	0.97	0.97	0.97	7	0.98	0.98	0.98	7	0.99	0.99	0.99	1

time: 32 ms (started: 2022-12-29 01:20:48 +05:00)

Analysis:

The audit risk dataset is a simple dataset therefore it has particularly good metrics. Since the dataset is balanced, I would focus on looking at accuracy for comparison:

1. The best accuracy achieved is 1.0 without applying any DR techniques. The total dims are 25.
2. After applying PCA, the accuracy slightly decreased to around 0.98 while reducing features significantly to 7 dimensions.
3. After trying other PCA variants, the results were at par with PCA with accuracy around 0.99.
4. After applying LDA, the best accuracy achieved was 0.95 which is slightly less than 1 however the number of dimensions become 1 which is a significant achievement.
5. SVD gives an accuracy of 0.98 and reduces dimensions to 7 which is the same as PCA.

All techniques perform quite well on this dataset. LDA compromises slightly on variance capture however it reduces dimensionality the most.

5.3 Regression Datasets

5.3.1 Loading Datasets

In [269...

```
power_df = pd.read_csv('Regression/CombinedCyclePowerPlantUCI.csv') #Combined Cycle Power Plant dataset
energy_df = pd.read_excel('Regression/EnergyEfficiencyUCI.xlsx') #Energy Efficiency Dataset
aquatic_df = pd.read_csv('Regression/qsar_aquatic_toxicityUCI.csv', sep = ';', names = ['TPSA', 'SAacc', 'H-050', 'MLC', 'RDCHI', 'GATS1p', 'nN', 'C-040', 'quantitative response_LC50'])

bikes_df = pd.read_csv('Regression/SeoulBikeDataUCI.csv', encoding='unicode_escape') #Bike Sharing Dataset
redwine_df = pd.read_csv('Regression/RedWineQualityUCI.csv', sep = ';') #Red Wine Quality Dataset
student_df = pd.read_csv('Regression/student-porUCI.csv', sep = ';') #Student Performance dataset
hardware_df = pd.read_csv('Regression/TomsHardwareUCI.data', names = ['NCD_0', 'NCD_1', 'NCD_2', 'NCD_3', 'NCD_4', 'NCD_5', 'NCD_6', 'NCD_7', 'NCD_8', 'NCD_9', 'BL_0', 'BL_1', 'BL_2', 'BL_3', 'BL_4', 'BL_5', 'BL_6', 'BL_7', 'BL_8', 'BL_9', 'NAD_0', 'NAD_1', 'NAD_2', 'NAD_3', 'NAD_4', 'NAD_5', 'NAD_6', 'NAD_7', 'NAD_8', 'NAD_9', 'AI_0', 'AI_1', 'AI_2', 'AI_3', 'AI_4', 'AI_5', 'AI_6', 'AI_7', 'AI_8', 'AI_9', 'NAC_0', 'NAC_1', 'NAC_2', 'NAC_3', 'NAC_4', 'NAC_5', 'NAC_6', 'NAC_7', 'NAC_8', 'NAC_9', 'ND_0', 'ND_1', 'ND_2', 'ND_3', 'ND_4', 'ND_5', 'ND_6', 'ND_7', 'ND_8', 'ND_9', 'CS_0', 'CS_1', 'CS_2', 'CS_3', 'CS_4', 'CS_5', 'CS_6', 'CS_7', 'CS_8', 'CS_9', 'AT_0', 'AT_1', 'AT_2', 'AT_3', 'AT_4', 'AT_5', 'AT_6', 'AT_7', 'AT_8', 'AT_9', 'NA_0', 'NA_1', 'NA_2', 'NA_3', 'NA_4', 'NA_5', 'NA_6', 'NA_7', 'NA_8', 'NA_9', 'ADL_0', 'ADL_1', 'ADL_2', 'ADL_3', 'ADL_4', 'ADL_5', 'ADL_6', 'ADL_7', 'ADL_8', 'ADL_9', 'AS(NA)_0', 'AS(NA)_1', 'AS(NA)_2', 'AS(NA)_3', 'AS(NA)_4', 'AS(NA)_5', 'AS(NA)_6', 'AS(NA)_7', 'AS(NA)_8', 'AS(NA)_9', 'AS(NAC)_0', 'AS(NAC)_1', 'AS(NAC)_2', 'AS(NAC)_3', 'AS(NAC)_4', 'AS(NAC)_5', 'AS(NAC)_6', 'AS(NAC)_7', 'AS(NAC)_8', 'AS(NAC)_9']) #Tom's Hardware Dataset

print('Shape of Combined Power Plant dataframe is: ', power_df.shape)
print('Shape of Energy Efficiency dataframe is: ', energy_df.shape)
print('Shape of Aquatic Toxicity dataframe is: ', aquatic_df.shape)
```

```
print('Shape of Seoul bike dataframe is: ', bikes_df.shape)
print('Shape of Red Wine Quality dataframe is: ', redwine_df.shape)
print('Shape of student dataframe is: ', student_df.shape)
print('Shape of Toms Hardware dataframe is: ', hardware_df.shape)
```

```
Shape of Combined Power Plant dataframe is: (9568, 5)
Shape of Energy Efficiency dataframe is: (768, 10)
Shape of Aquatic Toxicity dataframe is: (546, 8)
Shape of Seoul bike dataframe is: (8760, 13)
Shape of Red Wine Quality dataframe is: (1599, 12)
Shape of student dataframe is: (649, 33)
Shape of Toms Hardware dataframe is: (28179, 97)
time: 281 ms (started: 2022-12-26 16:14:57 +05:00)
```

5.3.2 Pre-Processing Datasets

1. Power Consumption Dataset

In [254...

```
print(power_df.isnull().sum()) #Check missing values
print(power_df.dtypes) #Check data types
power_df.head() #Bird's eye view of dataset
```

```
AT      0
V      0
AP      0
RH      0
PE      0
dtype: int64
AT      float64
V      float64
AP      float64
RH      float64
PE      float64
dtype: object
```

Out[254...

	AT	V	AP	RH	PE
0	8.34	40.77	1010.84	90.01	480.48
1	23.64	58.49	1011.40	74.20	445.75
2	29.74	56.90	1007.15	41.91	438.76
3	19.07	49.69	1007.22	76.79	453.09
4	11.80	40.66	1017.13	97.20	464.43

```
time: 0 ns (started: 2022-12-26 15:55:54 +05:00)
```

In [255...

```
power_classes = power_df[['PE']]
power_df.drop(columns = ['PE'], inplace = True) #Dropping unnecessary columns

#Scaling and One-hot encoding
Scaler = MinMaxScaler()
power_df = pd.get_dummies(power_df)
power_df = pd.DataFrame(Scaler.fit_transform(power_df), columns = power_df.columns)
print('Shape of df now is: ', power_df.shape)
power_df.head()
```

```
Shape of df now is: (9568, 4)
```

Out[255...

	AT	V	AP	RH
0	0.18	0.27	0.44	0.86
1	0.62	0.59	0.46	0.65
2	0.79	0.56	0.35	0.22
3	0.49	0.43	0.35	0.69
4	0.28	0.27	0.60	0.96

```
time: 16 ms (started: 2022-12-26 15:55:55 +05:00)
```

2. Energy Efficiency Dataset

In [256...

```
print(energy_df.isnull().sum()) #Check missing values
print(energy_df.dtypes) #Check data types
energy_df.head() #Bird's eye view of dataset
```

```
X1      0
X2      0
X3      0
X4      0
X5      0
```

```

X6      0
X7      0
X8      0
Y1      0
Y2      0
dtype: int64
X1      float64
X2      float64
X3      float64
X4      float64
X5      float64
X6      int64
X7      float64
X8      int64
Y1      float64
Y2      float64
dtype: object

```

```

Out[256...]

```

	X1	X2	X3	X4	X5	X6	X7	X8	Y1	Y2
0	0.98	514.50	294.00	110.25	7.00	2	0.00	0	15.55	21.33
1	0.98	514.50	294.00	110.25	7.00	3	0.00	0	15.55	21.33
2	0.98	514.50	294.00	110.25	7.00	4	0.00	0	15.55	21.33
3	0.98	514.50	294.00	110.25	7.00	5	0.00	0	15.55	21.33
4	0.90	563.50	318.50	122.50	7.00	2	0.00	0	20.84	28.28

time: 16 ms (started: 2022-12-26 15:55:56 +05:00)

```

In [257...]
energy_classes = energy_df[['Y1']]
energy_df.drop(columns = ['Y1', 'Y2'], inplace = True) #Dropping unnecessary columns

#Scaling and One-hot encoding
Scaler = MinMaxScaler()
energy_df = pd.get_dummies(energy_df)
energy_df = pd.DataFrame(Scaler.fit_transform(energy_df), columns = energy_df.columns)
print('Shape of df now is: ', energy_df.shape)
energy_df.head()

```

Shape of df now is: (768, 8)

```

Out[257...]

```

	X1	X2	X3	X4	X5	X6	X7	X8
0	1.00	0.00	0.29	0.00	1.00	0.00	0.00	0.00
1	1.00	0.00	0.29	0.00	1.00	0.33	0.00	0.00
2	1.00	0.00	0.29	0.00	1.00	0.67	0.00	0.00
3	1.00	0.00	0.29	0.00	1.00	1.00	0.00	0.00
4	0.78	0.17	0.43	0.11	1.00	0.00	0.00	0.00

time: 16 ms (started: 2022-12-26 15:55:56 +05:00)

3. Aquatic Toxicity Dataset

```

In [258...]
print(aquatic_df.isnull().sum()) #Check missing values
print(aquatic_df.dtypes) #Check data types
aquatic_df.head() #Bird's eye view of dataset

```

```

TPSA      0
SAacc      0
H-050      0
MLOGPRDCHI 0
GATS1p      0
nN          0
C-040      0
quantitative response_LC50 0
dtype: int64
TPSA      float64
SAacc      int64
H-050      float64
MLOGPRDCHI float64
GATS1p      float64
nN          int64
C-040      int64
quantitative response_LC50 float64
dtype: object

```

```

Out[258...]

```

	TPSA	SAacc	H-050	MLOGPRDCHI	GATS1p	nN	C-040	quantitative response_LC50
0.00	0.00	0	2.42		1.23	0.67	0	3.74
0.00	0.00	0	2.64		1.40	0.63	0	4.33

	TPSA	SAacc	H-050	MLOGPRDCHI	GATS1p	nN	C-040	quantitative response_LC50
9.23	11.00	0	5.80	2.93	0.49	0	0	7.02
9.23	11.00	0	5.45	2.89	0.49	0	0	6.72
9.23	11.00	0	4.07	2.76	0.69	0	0	5.98

time: 15 ms (started: 2022-12-26 15:56:12 +05:00)

In [259...

```

aquatic_classes = aquatic_df[['quantitative response_LC50']]
aquatic_df.drop(columns = ['quantitative response_LC50'], inplace = True) #Dropping unnecessary columns

#Scaling and One-hot encoding
Scaler = MinMaxScaler()
aquatic_df = pd.get_dummies(aquatic_df)
aquatic_df = pd.DataFrame(Scaler.fit_transform(aquatic_df), columns = aquatic_df.columns)
print('Shape of df now is: ', aquatic_df.shape)
aquatic_df.head()

```

Shape of df now is: (546, 7)

Out[259...

	TPSA	SAacc	H-050	MLOGPRDCHI	GATS1p	nN	C-040
0	0.00	0.00	0.57	0.04	0.17	0.00	0.00
1	0.00	0.00	0.58	0.07	0.16	0.00	0.00
2	0.02	0.00	0.79	0.35	0.09	0.00	0.00
3	0.02	0.00	0.76	0.35	0.10	0.00	0.00
4	0.02	0.00	0.67	0.32	0.19	0.00	0.00

time: 0 ns (started: 2022-12-26 15:56:43 +05:00)

4. Seoul Bikes Dataset

In [195...

```

print(bikes_df.isnull().sum()) #Check missing values
print(bikes_df.dtypes) #Check data types
bikes_df.head() #Bird's eye view of dataset

```

```

Rented Bike Count      0
Hour                   0
Temperature(°C)         0
Humidity(%)             0
Wind speed (m/s)        0
Visibility (10m)         0
Dew point temperature(°C) 0
Solar Radiation (MJ/m2) 0
Rainfall(mm)            0
Snowfall (cm)           0
Seasons                 0
Holiday                 0
Functioning Day         0
dtype: int64
Rented Bike Count      int64
Hour                   int64
Temperature(°C)        float64
Humidity(%)            int64
Wind speed (m/s)        float64
Visibility (10m)        int64
Dew point temperature(°C) float64
Solar Radiation (MJ/m2) float64
Rainfall(mm)           float64
Snowfall (cm)          float64
Seasons                object
Holiday                object
Functioning Day         object
dtype: object

```

Out[195...

	Rented Bike Count	Hour	Temperature(°C)	Humidity(%)	Wind speed (m/s)	Visibility (10m)	Dew point temperature(°C)	Solar Radiation (MJ/m2)	Rainfall(mm)	Snowfall (cm)	Seasons	Holiday	Fu
0	254	0	-5.20	37	2.20	2000	-17.60	0.00	0.00	0.00	Winter	No Holiday	
1	204	1	-5.50	38	0.80	2000	-17.60	0.00	0.00	0.00	Winter	No Holiday	
2	173	2	-6.00	39	1.00	2000	-17.70	0.00	0.00	0.00	Winter	No Holiday	
3	107	3	-6.20	40	0.90	2000	-17.60	0.00	0.00	0.00	Winter	No Holiday	
4	78	4	-6.00	36	2.30	2000	-18.60	0.00	0.00	0.00	Winter	No	

Rented Bike Count	Hour	Temperature(°C)	Humidity(%)	Wind speed (m/s)	Visibility (10m)	Dew point temperature(°C)	Solar Radiation (MJ/m2)	Rainfall(mm)	Snowfall (cm)	Seasons	Holiday	Fu
Holiday												

time: 16 ms (started: 2022-12-26 03:34:54 +05:00)

In [196...

```
bikes_classes = bikes_df[['Rented Bike Count']]
bikes_df.drop(columns = ['Rented Bike Count'], inplace = True) #Dropping unnecessary columns

#Scaling and One-hot encoding
Scaler = MinMaxScaler()
bikes_df = pd.get_dummies(bikes_df)
bikes_df = pd.DataFrame(Scaler.fit_transform(bikes_df), columns = bikes_df.columns)
print('Shape of df now is: ', bikes_df.shape)
bikes_df.head()
```

Shape of df now is: (8760, 17)

Out[196...

	Hour	Temperature(°C)	Humidity(%)	Wind speed (m/s)	Visibility (10m)	Dew point temperature(°C)	Solar Radiation (MJ/m2)	Rainfall(mm)	Snowfall (cm)	Seasons_Autumn	Seasons_Spr
0	0.00	0.22	0.38	0.30	1.00	0.22	0.00	0.00	0.00	0.00	(
1	0.04	0.22	0.39	0.11	1.00	0.22	0.00	0.00	0.00	0.00	(
2	0.09	0.21	0.40	0.14	1.00	0.22	0.00	0.00	0.00	0.00	(
3	0.13	0.20	0.41	0.12	1.00	0.22	0.00	0.00	0.00	0.00	(
4	0.17	0.21	0.37	0.31	1.00	0.21	0.00	0.00	0.00	0.00	(

time: 16 ms (started: 2022-12-26 03:34:54 +05:00)

5. Red Wine Quality Dataset

In [197...

```
print(redwine_df.isnull().sum()) #Check missing values
print(redwine_df.dtypes) #Check data types
redwine_df.head() #Bird's eye view of dataset
```

```
fixed acidity      0
volatile acidity   0
citric acid        0
residual sugar     0
chlorides          0
free sulfur dioxide 0
total sulfur dioxide 0
density           0
pH               0
sulphates         0
alcohol           0
quality           0
dtype: int64
fixed acidity      float64
volatile acidity   float64
citric acid        float64
residual sugar     float64
chlorides          float64
free sulfur dioxide float64
total sulfur dioxide float64
density           float64
pH               float64
sulphates         float64
alcohol           float64
quality           int64
dtype: object
```

Out[197...

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.40	0.70	0.00	1.90	0.08	11.00	34.00	1.00	3.51	0.56	9.40	5
1	7.80	0.88	0.00	2.60	0.10	25.00	67.00	1.00	3.20	0.68	9.80	5
2	7.80	0.76	0.04	2.30	0.09	15.00	54.00	1.00	3.26	0.65	9.80	5
3	11.20	0.28	0.56	1.90	0.07	17.00	60.00	1.00	3.16	0.58	9.80	6
4	7.40	0.70	0.00	1.90	0.08	11.00	34.00	1.00	3.51	0.56	9.40	5

time: 0 ns (started: 2022-12-26 03:34:54 +05:00)

In [198...

```
redwine_classes = redwine_df[['quality']]
redwine_df.drop(columns = ['quality'], inplace = True) #Dropping unnecessary columns

#Scaling and One-hot encoding
Scaler = MinMaxScaler()
redwine_df = pd.get_dummies(redwine_df)
redwine_df = pd.DataFrame(Scaler.fit_transform(redwine_df), columns = redwine_df.columns)
print('Shape of df now is: ', redwine_df.shape)
redwine_df.head()
```

Shape of df now is: (1599, 11)

Out[198...

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	0.25	0.40	0.00	0.07	0.11	0.14	0.10	0.57	0.61	0.14	0.15
1	0.28	0.52	0.00	0.12	0.14	0.34	0.22	0.49	0.36	0.21	0.22
2	0.28	0.44	0.04	0.10	0.13	0.20	0.17	0.51	0.41	0.19	0.22
3	0.58	0.11	0.56	0.07	0.11	0.23	0.19	0.58	0.33	0.15	0.22
4	0.25	0.40	0.00	0.07	0.11	0.14	0.10	0.57	0.61	0.14	0.15

time: 16 ms (started: 2022-12-26 03:34:55 +05:00)

6. Student Portugese Dataset

In [270...

```
print(student_df.isnull().sum()) #Check missing values
print(student_df.dtypes) #Check data types
student_df.head() #Bird's eye view of dataset
```

```
school      0
sex         0
age         0
address     0
famsize     0
Pstatus     0
Medu        0
Fedu        0
Mjob        0
Fjob        0
reason      0
guardian    0
traveltime  0
studytime   0
failures    0
schoolsup   0
famsup       0
paid        0
activities  0
nursery     0
higher      0
internet    0
romantic    0
famrel      0
freetime    0
goout       0
Dalc        0
Walc        0
health      0
absences    0
G1          0
G2          0
G3          0
dtype: int64
school      object
sex         object
age         int64
address     object
famsize     object
Pstatus     object
Medu        int64
Fedu        int64
Mjob        object
Fjob        object
reason      object
guardian    object
traveltime  int64
studytime   int64
failures    int64
schoolsup   object
famsup       object
paid        object
```

```

activities    object
nursery       object
higher        object
internet      object
romantic      object
famrel        int64
freetime      int64
goout         int64
Dalc          int64
Walc          int64
health        int64
absences      int64
G1            int64
G2            int64
G3            int64
dtype: object

```

```

Out[270...]
   school sex  age address famsize Pstatus Medu Fedu  Mjob  Fjob ... famrel freetime goout Dalc Walc health absences
0      GP  F   18      U    GT3      A      4      4  at_home teacher ...      4      3      4      1      1      3      4
1      GP  F   17      U    GT3      T      1      1  at_home  other ...      5      3      3      1      1      3      2
2      GP  F   15      U    LE3      T      1      1  at_home  other ...      4      3      2      2      3      3      6
3      GP  F   15      U    GT3      T      4      2  health  services ...      3      2      2      1      1      5      0
4      GP  F   16      U    GT3      T      3      3   other   other ...      4      3      2      1      2      5      0

```

5 rows × 33 columns

time: 0 ns (started: 2022-12-26 16:15:05 +05:00)

```

In [271...]
student_classes = student_df[['G3']]
student_df.drop(columns = [ 'G3'], inplace = True) #Dropping unnecessary columns
Scaler = MinMaxScaler()

#Scaling and One-hot encoding
student_df = pd.get_dummies(student_df)
student_df = pd.DataFrame(Scaler.fit_transform(student_df), columns = student_df.columns)
print('Shape of df now is: ', student_df.shape)
student_df.head()

```

Shape of df now is: (649, 58)

```

Out[271...]
   age  Medu  Fedu  traveltime  studytime  failures  famrel  freetime  goout  Dalc ... activities_no  activities_yes  nursery_no  nursery_yes
0  0.43   1.00   1.00      0.33      0.33      0.00   0.75      0.50   0.75  0.00 ...          1.00          0.00          0.00          1.00
1  0.29   0.25   0.25      0.00      0.33      0.00   1.00      0.50   0.50  0.00 ...          1.00          0.00          1.00          0.00
2  0.00   0.25   0.25      0.00      0.33      0.00   0.75      0.50   0.25  0.25 ...          1.00          0.00          0.00          1.00
3  0.00   1.00   0.50      0.00      0.67      0.00   0.50      0.25   0.25  0.00 ...          0.00          1.00          0.00          1.00
4  0.14   0.75   0.75      0.00      0.33      0.00   0.75      0.50   0.25  0.00 ...          1.00          0.00          0.00          1.00

```

5 rows × 58 columns

time: 16 ms (started: 2022-12-26 16:15:14 +05:00)

7. Tom's Hardware Dataset

```

In [201...]
print(hardware_df.isnull().sum()) #Check missing values
print(hardware_df.dtypes)        #Check data types
hardware_df.head()               #Bird's eye view of dataset

```

```

NCD_0      0
NCD_1      0
NCD_2      0
NCD_3      0
NCD_4      0
..
AS(NAC)_3  0
AS(NAC)_4  0
AS(NAC)_5  0
AS(NAC)_6  0
AS(NAC)_7  0
Length: 97, dtype: int64
NCD_0      int64
NCD_1      int64
NCD_2      int64
NCD_3      int64
NCD_4      int64
...

```



```
AS(NAC)_3    float64
AS(NAC)_4    float64
AS(NAC)_5    float64
AS(NAC)_6    float64
AS(NAC)_7    float64
Length: 97, dtype: object
```

```
Out[201...]
   NCD_0  NCD_1  NCD_2  NCD_3  NCD_4  NCD_5  NCD_6  NCD_7  BL_0  BL_1  ...  AS(NA)_6  AS(NA)_7  AS(NAC)_0  AS(NAC)_1  AS(NAC)_
0      0      0      0      0      0      0      0      0      0  0.00  0.00  ...      0.00      0.00      0.00      0.00      0.0
1      0      0      0      0      0      0      0      0      0  0.00  0.00  ...      0.00      0.00      0.00      0.00      0.0
2      0      0      0      0      0      0      0      0      0  0.00  0.00  ...      0.00      0.00      0.00      0.00      0.0
3      0      0      0      0      0      0      0      0      0  0.00  0.00  ...      0.00      0.00      0.00      0.00      0.0
4      0      0      0      0      0      0      0      0      0  0.00  0.00  ...      0.00      0.00      0.00      0.00      0.0
```

5 rows × 97 columns

time: 15 ms (started: 2022-12-26 03:34:58 +05:00)

```
In [202...]
hardware_classes = hardware_df[['AS(NAC)_7']]
hardware_df.drop(columns = ['AS(NAC)_7'], inplace = True) #Dropping unnecessary columns

#Scaling and One-hot encoding
Scaler = MinMaxScaler()
hardware_df = pd.get_dummies(hardware_df)
hardware_df = pd.DataFrame(Scaler.fit_transform(hardware_df), columns = hardware_df.columns)
print('Shape of df now is: ', hardware_df.shape)
hardware_df.head()
```

Shape of df now is: (28179, 96)

```
Out[202...]
   NCD_0  NCD_1  NCD_2  NCD_3  NCD_4  NCD_5  NCD_6  NCD_7  BL_0  BL_1  ...  AS(NA)_5  AS(NA)_6  AS(NA)_7  AS(NAC)_0  AS(NAC)_1
0  0.00  0.00  0.00  0.00  0.00  0.00  0.00  0.00  0.00  0.00  ...      0.00      0.00      0.00      0.00      0.00
1  0.00  0.00  0.00  0.00  0.00  0.00  0.00  0.00  0.00  0.00  ...      0.00      0.00      0.00      0.00      0.00
2  0.00  0.00  0.00  0.00  0.00  0.00  0.00  0.00  0.00  0.00  ...      0.00      0.00      0.00      0.00      0.00
3  0.00  0.00  0.00  0.00  0.00  0.00  0.00  0.00  0.00  0.00  ...      0.00      0.00      0.00      0.00      0.00
4  0.00  0.00  0.00  0.00  0.00  0.00  0.00  0.00  0.00  0.00  ...      0.00      0.00      0.00      0.00      0.00
```

5 rows × 96 columns

time: 62 ms (started: 2022-12-26 03:35:10 +05:00)

5.3.3 Dimensionality Reduction Techniques with ML

A) Dimensionality Reduction Pipeline for Regression Datasets:

```
In [221...]
def dimensionality_reduction_regression(X_train, y_train, X_test, y_test, pca_dim = 0.95, ipca_dim= 0.95, svd_dim = 0

    print("Starting DR Pipeline...")

    print("1. Running Lazy Predict without DR")

    #Running Lazy Predict without Dimensionality Reduction:
    clf = LazyRegressor(verbose=0, ignore_warnings=True, custom_metric = None)
    simple_models = clf.fit(X_train, X_test, y_train, y_test)[0].sort_index()
    simple_models['dim'] = X_train.shape[1]
    simple_models.drop(columns = ['Time Taken'], inplace = True)

    print("Success!")

    #Model Performances with PCA:

    #Running PCA:
    print("2. Running PCA")
    pca = PCA(n_components = pca_dim)
    pca.fit(X_train)
    X_train_transformed = pca.transform(X_train)
    X_test_transformed = pca.transform(X_test)
    print("Success!")

    #Running Lazy Predict with PCA
    print("3. Running Lazy Predict on PCA dataset")
    clf = LazyRegressor(verbose=0, ignore_warnings=True, custom_metric = None)
    pcamodels = clf.fit(X_train_transformed, X_test_transformed, y_train, y_test)[0].sort_index()
    pcamodels.drop(columns = ['Time Taken'], inplace = True)
```

```
pcamodels['dims'] = len(pca.components_)
print("Success!")
```

#Model Performances with Incremental-PCA:

```
#Running Incremental PCA:
print("4. Running Incremental PCA")
for i in range(1, X_train.shape[1], 1):
    ipca = IncrementalPCA(n_components = i)
    ipca.fit(X_train)
    X_train_transformed = ipca.transform(X_train)
    X_test_transformed = ipca.transform(X_test)

    if ipca.explained_variance_ratio_.sum() >= ipca_dim:
        print("Success!")
        print("5. Running Lazy Predict on Incremental PCA dataset")
        clf = LazyRegressor(verbose=0, ignore_warnings=True, custom_metric = None) #Running Lazy Predict after
        ipcamodels = clf.fit(X_train_transformed, X_test_transformed, y_train, y_test)[0].sort_index()
        ipcamodels.drop(columns = [ 'Time Taken'], inplace = True)
        ipcamodels['dims'] = ipca.n_components_
        print("Success!")
        break
```

#Model Performances with Sparse-PCA:

```
print("6. Running Sparse PCA")
spca = SparsePCA(n_components = 10)
spca.fit(X_train)
X_train_transformed = spca.transform(X_train)
X_test_transformed = spca.transform(X_test)
print("Success!")
```

#Running Lazy Predict Sparse PCA:

```
print("7. Running Lazy Predict on Sparse PCA dataset")
clf = LazyRegressor(verbose=0, ignore_warnings=True, custom_metric = None) #Running Lazy Predict after SVD
spcamodels = clf.fit(X_train_transformed, X_test_transformed, y_train, y_test)[0].sort_index()
spcamodels.drop(columns = [ 'Time Taken'], inplace = True)
spcamodels['dims'] = spca.n_components_
print("Success!")
```

#Model Performances with SVD:

```
#Running SVD:
print("8. Running SVD")
for i in range(1, X_train.shape[1], 1):
    svd = TruncatedSVD(n_components = i)
    svd.fit(X_train)
    X_train_transformed = svd.transform(X_train)
    X_test_transformed = svd.transform(X_test)

    if svd.explained_variance_ratio_.sum() >= svd_dim or i>=X_train.shape[1]-1:
        print("Success!")
        print("9. Running Lazy Predict on SVD dataset")
        clf = LazyRegressor(verbose=0, ignore_warnings=True, custom_metric = None) #Running Lazy Predict after
        svdmodels = clf.fit(X_train_transformed, X_test_transformed, y_train, y_test)[0].sort_index()
        svdmodels.drop(columns = [ 'Time Taken'], inplace = True)
        svdmodels['dims'] = len(svd.components_)
        print("Success!")
        break
```

#Compiling Model Results:

```
print("Compiling Model Results")
models_results = pd.concat([simple_models,
                             pcamodels,
                             ipcamodels,
                             spcamodels,
                             svdmodels], axis = 1, keys = ['Without DR', 'PCA ', 'Incremental-PCA', 'Sparse-PCA', 'SVD'])

print("Pipeline run Successful")

return models_results
```

time: 0 ns (started: 2022-12-26 03:51:55 +05:00)

B) Applying Pipeline to Regression Datasets

1. Power Consumption Dataset

In [226...

```
#Lets split the data into a train test split from the start, test set will be kept separate and will only be used for

X_train, X_test, y_train, y_test = train_test_split(power_df, power_classes, test_size=0.25, random_state =43)
```

In [227...

In [337...

Out[337...]

1. When no DR technique is applied, the best adjusted R2 score achieved was 0.96.

2. After applying PCA, the R2 score slightly decreased to 0.93 with a reduction in features to 3 from 4.
3. When applying other PCA variants, Incremental PCA behaved like PCA, but sparse PCA captured all the variance however it increased the dimensionality rather than decreasing it.
4. When SVD was tried, the variance captured was 0.93 slightly less than 0.96 with 1 feature reduction so it performed at par with PCA.

Only 1 feature was reduced in these DR techniques with little compromise to variance capture, PCA and SVD performed at par with each other!

2. Energy Efficiency Dataset

In [228...

```
#Lets split the data into a train test split from the start, test set will be kept separate and will only be used for
```

```
X_train, X_test, y_train, y_test = train_test_split(energy_df, energy_classes, test_size=0.25, random_state =43)
```

time: 0 ns (started: 2022-12-26 09:38:36 +05:00)

In [229...

```
#Pipeline run
```

```
models_results = dimensionality_reduction_regression(X_train, y_train, X_test, y_test, pca_dim = 0.95, svd_dim = 0.95)
```

```
#keeping only the desired algorithms
```

```
results = models_results.T[['AdaBoostRegressor', 'DecisionTreeRegressor',
```

```
                            'ElasticNetCV', 'GradientBoostingRegressor', 'KNeighborsRegressor', 'XGBRegressor',
```

```
                            'RandomForestRegressor', 'SVR']].T
```

```
#Exporting results to excel
```

```
results.to_excel('Results/Regression/Energy.xlsx', sheet_name = 'Energy Efficiency Dataset')
```

Starting DR Pipeline...

1. Running Lazy Predict without DR

100%

42/42 [00:02<00:00, 14.14it/s]

Success!

2. Running PCA

Success!

3. Running Lazy Predict on PCA dataset

100%

42/42 [00:02<00:00, 14.01it/s]

Success!

4. Running Incremental PCA

Success!

5. Running Lazy Predict on Incremental PCA dataset

100%

42/42 [00:02<00:00, 14.03it/s]

Success!

6. Running Sparse PCA

Success!

7. Running Lazy Predict on Sparse PCA dataset

100%

42/42 [00:03<00:00, 13.78it/s]

Success!

8. Running SVD

Success!

9. Running Lazy Predict on SVD dataset

100%

42/42 [00:03<00:00, 13.33it/s]

Success!

Compiling Model Results

Pipeline run Successful

time: 15.4 s (started: 2022-12-26 09:38:36 +05:00)

Dataset Results:

In [338...

```
#Results
```

```
results = pd.read_excel('Results/Regression/Energy.xlsx', header=[0, 1], index_col=0 )
```

```
results
```

Out[338...

Model	Without DR				PCA				Incremental-PCA				S
	Adjusted R-Squared	R-Squared	RMSE	dim	Adjusted R-Squared	R-Squared	RMSE	dims	Adjusted R-Squared	R-Squared	RMSE	dims	
AdaBoostRegressor	0.97	0.97	1.84	8	0.95	0.95	2.32	5	0.95	0.95	2.23	5	0.97
DecisionTreeRegressor	1.00	1.00	0.47	8	0.98	0.98	1.41	5	0.98	0.98	1.42	5	1.00

	Model	Without DR				PCA				Incremental-PCA				Size
		Adjusted R-Squared	R-Squared	RMSE	dim	Adjusted R-Squared	R-Squared	RMSE	dims	Adjusted R-Squared	R-Squared	RMSE	dims	
Model	ElasticNetCV	0.93	0.93	2.68	8	0.92	0.92	2.88	5	0.92	0.92	2.88	5	0.93
	GradientBoostingRegressor	1.00	1.00	0.45	8	0.99	0.99	0.96	5	0.99	0.99	0.93	5	1.00
	KNeighborsRegressor	0.96	0.96	2.01	8	0.94	0.94	2.50	5	0.94	0.94	2.50	5	0.96
	XGBRegressor	1.00	1.00	0.29	8	1.00	1.00	0.70	5	0.99	0.99	0.72	5	1.00
	RandomForestRegressor	1.00	1.00	0.44	8	0.99	0.99	0.76	5	0.99	0.99	0.78	5	1.00
	SVR	0.94	0.94	2.51	8	0.91	0.91	2.99	5	0.91	0.91	2.99	5	0.93

Analysis:

1. When no DR technique was applied, the best Adjusted R2 score achieved was 1 which means the model captures all the variance available.
2. When PCA is applied, the features reduced to 5 from 8 but the best adjusted R2 score remained at 1 which shows PCA to be a good technique for this dataset.
3. As other PCA variants are tried, incremental PCA performs at par with PCA, but Sparse PCA increases dimension, so it performs poorly.
4. SVD reduced the dimensions dataset dimensions to 6 from 8 only with almost no compromise on variance capture.

3. Aquative Toxicity Dataset

```
X_train, X_test, y_train, y_test = train_test_split(aquatic_df, aquatic_classes, test_size=0.25, random_state =43)
```

In [261...

```
#Pipeline run
models_results = dimensionality_reduction_regression(X_train, y_train, X_test, y_test, pca_dim = 0.95, svd_dim = 0.95)

#keeping only the desired algorithms
results = models_results.T[['AdaBoostRegressor', 'DecisionTreeRegressor',
                           'ElasticNetCV', 'GradientBoostingRegressor', 'KNeighborsRegressor', 'XGBRegressor',
                           'RandomForestRegressor', 'SVR']].T

#Exporting results to excel
results.to_excel('Results/Regression/aquatic.xlsx', sheet_name = 'Aquatic Toxicity Dataset')
```

- ### 1. Running Lazy Predict without DR

Success!

Success!

```
s] |██████████| 42/42 [00:01<00:00, 22.72it/
```

Success!

Success!

[illegible]

Success!

Success!

```
s] 100%|███████████ 42/42 [00:01<00:00, 21.07it/
```

Success!

8. Running SVD

Success!

9. Running Lazy Predict on SVD dataset

Dataset Results:

```
#Results
results = pd.read_excel('Results/Regression/Aquatic.xlsx', header=[0, 1], index_col=0)
results
```

Model	Without DR					PCA				Incremental-PCA				Size
	Adjusted R-Squared	R-Squared	RMSE	dim	Adjusted R-Squared	R-Squared	RMSE	dims	Adjusted R-Squared	R-Squared	RMSE	dims	Adjusted R-Squared	
AdaBoostRegressor	0.36	0.39	1.22	7	0.29	0.32	1.30	5	0.35	0.37	1.24	5	0.44	
DecisionTreeRegressor	0.26	0.30	1.32	7	-0.10	-0.06	1.61	5	0.03	0.06	1.52	5	0.24	
ElasticNetCV	0.39	0.42	1.20	7	0.40	0.42	1.19	5	0.40	0.42	1.19	5	0.37	
GradientBoostingRegressor	0.42	0.45	1.17	7	0.34	0.37	1.25	5	0.40	0.42	1.19	5	0.40	
KNeighborsRegressor	0.36	0.40	1.22	7	0.41	0.43	1.18	5	0.40	0.43	1.19	5	0.35	
XGBRegressor	0.37	0.40	1.21	7	0.26	0.29	1.32	5	0.27	0.30	1.31	5	0.36	
RandomForestRegressor	0.42	0.45	1.16	7	0.42	0.44	1.17	5	0.43	0.45	1.16	5	0.42	
SVR	0.42	0.45	1.16	7	0.41	0.43	1.18	5	0.41	0.43	1.18	5	0.41	

1. The best adjusted R2 score achieved without any DR technique was 0.42.
2. After applying PCA, the number of features reduced from 7 to 5 and adjusted R2 improved to 0.44.
3. When other variants of PCA were tried, incremental PCA performed at par with PCA whereas sparse PCA increased dimensions instead.
4. When SVD is applied, the performance improves further to 0.47 whereas the number of features reduces from 7 to 5.

4. Seoul Bikes Dataset

```
X_train, X_test, y_train, y_test = train_test_split(bikes_df, bikes_classes, test_size=0.25, random_state=43)
```

```
#Pipeline run
models_results = dimensionality_reduction_regression(X_train, y_train, X_test, y_test, pca_dim = 0.95, svd_dim = 0.95)

#keeping only the desired algorithms
results = models_results.T[['AdaBoostRegressor', 'DecisionTreeRegressor',
                           'ElasticNetCV', 'GradientBoostingRegressor', 'KNeighborsRegressor', 'XGBRegressor',
                           'RandomForestRegressor', 'SVR']].T

#Exporting results to excel
results.to_excel('Results/Regression/Bikes.xlsx', sheet_name = 'Seoul Bikes Dataset')
```

4. Running Incremental PCA


```
Starting DR Pipeline...
1. Running Lazy Predict without DR
100%|██████████████████████████████████████████████████████████████████████████████| 42/42 [00:08<00:00, 5.13it/s]
Success!
2. Running PCA
Success!
3. Running Lazy Predict on PCA dataset
100%|██████████████████████████████████████████████████████████████████████████████| 42/42 [00:08<00:00, 5.11it/s]
Success!
4. Running Incremental PCA
Success!
5. Running Lazy Predict on Incremental PCA dataset
100%|██████████████████████████████████████████████████████████████████████████████| 42/42 [00:08<00:00, 5.17it/s]
Success!
6. Running Sparse PCA
Success!
7. Running Lazy Predict on Sparse PCA dataset
100%|██████████████████████████████████████████████████████████████████████████████| 42/42 [00:08<00:00, 5.20it/s]
Success!
8. Running SVD
Success!
9. Running Lazy Predict on SVD dataset
100%|██████████████████████████████████████████████████████████████████████████████| 42/42 [00:08<00:00, 5.12it/s]
Success!
Compiling Model Results
Pipeline run Successful
time: 41.2 s (started: 2022-12-26 14:08:20 +05:00)
```

In [341...

Out[341...]

```
time: 15 ms (started: 2022-12-29 01:26:26 +05:00)
```

PCA, its variants and SVD all perform well for this dataset!

6. Student Portuguese Dataset

In [272...

```
X_train, X_test, y_train, y_test = train_test_split(student_df, student_classes, test_size=0.25, random_state =43)
```

```
time: 0 ns (started: 2022-12-26 16:15:29 +05:00)
```

In [273...

```
#Pipeline run
models_results = dimensionality_reduction_regression(X_train, y_train, X_test, y_test, pca_dim = 0.95, svd_dim = 0.95)

#keeping only the desired algorithms
results = models_results.T[['AdaBoostRegressor', 'DecisionTreeRegressor',
                            'ElasticNetCV', 'GradientBoostingRegressor', 'KNeighborsRegressor', 'XGBRegressor',
                            'RandomForestRegressor', 'SVR']].T

#Exporting results to excel
results.to_excel('Results/Regression/Student.xlsx', sheet_name = "Student Portugese Dataset")
```

Starting DR Pipeline...

1. Running Lazy Predict without DR

[illegible]

Success!

2. Running PCA

Success!

3. Running Lazy Predict on PCA dataset

[illegible]

Success!

4. Running Incremental PCA

Success!

5. Running Lazy Predict on Incremental PCA dataset

[illegible]

Success!

6. Running Sparse PCA

Success!

7. Running Lazy Predict on Sparse PCA dataset

[illegible]

Success!

8. Running SVD

Success!

9. Running Lazy Predict on SVD dataset

[illegible]

Success!

Compiling Model Results

Pipeline run Successful

```
time: 15.5 s (started: 2022-12-26 16:15:30 +05:00)
```

Dataset Results:

In [342...

```
#Results
results = pd.read_excel('Results/Regression/Student.xlsx', header=[0, 1], index_col=0)
results
```

Out[342...]

Model	Without DR					PCA				Incremental-PCA				Spearman Rank
	Adjusted R-Squared	R-Squared	RMSE	dim	Adjusted R-Squared	R-Squared	RMSE	dims	Adjusted R-Squared	R-Squared	RMSE	dims	Adjusted R-Squared	
AdaBoostRegressor	0.69	0.80	1.47	58	0.04	0.21	2.91	29	-0.02	0.16	3.00	29	-0.04	
DecisionTreeRegressor	0.46	0.65	1.93	58	-1.12	-0.74	4.32	29	-0.87	-0.54	4.06	29	-1.46	
ElasticNetCV	0.73	0.83	1.36	58	0.09	0.25	2.83	29	0.09	0.25	2.83	29	0.08	
GradientBoostingRegressor	0.69	0.80	1.46	58	0.03	0.21	2.92	29	0.03	0.20	2.93	29	-0.14	
KNeighborsRegressor	0.18	0.47	2.38	58	0.01	0.19	2.95	29	-0.06	0.13	3.05	29	-0.11	
XGBRegressor	0.64	0.77	1.57	58	-0.02	0.16	3.00	29	0.01	0.19	2.95	29	-0.33	
RandomForestRegressor	0.70	0.81	1.43	58	0.09	0.25	2.83	29	0.07	0.24	2.86	29	0.00	
SVR	0.49	0.67	1.87	58	0.10	0.26	2.82	29	0.08	0.24	2.85	29	0.02	

Analysis:

1. When no DR technique was applied, the best Adjusted R2 score achieved was 0.70.
2. When PCA is applied, the features reduced from 58 to 29 but the best the adjusted R2 score drops significantly which means PCA is not a good DR technique for this dataset.
3. The rest of the PCA variants and SVD performs poorly as well! **This dataset does not seem to allow any feature reduction which could be highlighting that there is a remarkably high correlation between some variables with predicted value!**

In [238...

time: 94 ms (started: 2022-12-26 14:09:39 +05:00)

In [239...

Starting DR Pipeline...

[illegible]

2. Running PCA

3. Running Lazy Predict on PCA dataset

[illegible]

4. Running Incremental PCA

5. Running Lazy Predict on Incremental PCA dataset

[illegible]

Success!

6. Running Sparse PCA

Success!

7. Running Lazy Predict on Sparse PCA dataset

[illegible]

Success!

8. Running SVD

Success!

9. Running Lazy Predict on SVD dataset

[illegible]

Success!

Compiling Model Results

Pipeline run Successful

```
time: 56min 41s (started: 2022-12-26 14:09:39 +05:00)
```

In [343...

Out[343]:

Model	Without DR				PCA				Incremental-PCA				
	Adjusted R-Squared	R-Squared	RMSE	dim	Adjusted R-Squared	R-Squared	RMSE	dims	Adjusted R-Squared	R-Squared	RMSE	dims	Adjusted R-Squared
AdaBoostRegressor	0.88	0.88	4458.21	96	-1.85	-1.85	21948.76	18	-1.61	-1.61	21005.85	18	-1.61
DecisionTreeRegressor	0.92	0.92	3697.34	96	0.79	0.79	5893.47	18	0.82	0.82	5523.79	18	0.82

Model	Without DR						PCA		Incremental-PCA				
	Adjusted R-Squared	R-Squared	RMSE	dim	Adjusted R-Squared	R-Squared	RMSE	dims	Adjusted R-Squared	R-Squared	RMSE	dims	Adjusted R-Squared
ElasticNetCV	0.83	0.83	5304.48	96	0.15	0.16	11946.31	18	0.15	0.16	11946.38	18	
GradientBoostingRegressor	0.97	0.97	2354.97	96	0.90	0.90	4118.33	18	0.90	0.90	4021.61	18	
KNeighborsRegressor	0.95	0.95	3006.69	96	0.88	0.88	4498.94	18	0.88	0.88	4495.64	18	
XGBRegressor	0.97	0.97	2290.19	96	0.90	0.90	4063.83	18	0.90	0.90	4057.10	18	
RandomForestRegressor	0.97	0.97	2297.70	96	0.90	0.90	4143.08	18	0.90	0.90	4199.52	18	
SVR	-0.03	-0.02	13124.97	96	-0.03	-0.03	13216.94	18	-0.03	-0.03	13216.98	18	

time: 15 ms (started: 2022-12-29 01:27:20 +05:00)

Analysis:

Adjusted R2 score is used to compare the results for the Tom's Hardware dataset:

1. When no DR technique was applied, the best Adjusted R2 score achieved was 0.97 which means the model captures almost all the variance available.
2. When PCA is applied, the features reduced from 96 to 18 but with a small drop in adjusted R2 score of 0.06 which leads to an R2 score of 0.90.
3. As other PCA variants are tried, incremental PCA performs at par with PCA but sparse reduces dimensions to 10 with some more compromise on adjusted R2 value leading to R2 score of 0.87
4. SVD performs at par with PCA as it leads to an adjusted R2 score of 0.9 and gives a feature set of 18 dimensions.

PCA and SVD both perform well as they reduce the number of features from 96 to 18 and capture all variance!

6. Critical Analysis

All the datasets do not compromise a lot on variance capture when their dimensionality is reduced using any of the techniques tried above which proves that these techniques are especially useful and should be implemented before diving deep into machine learning. Spending some time on reducing dimensions would help in the long run since model development becomes easier and less time-consuming when dataset features are reduced. Furthermore, LDA is only applicable on classification datasets and there are other varying factors which does not allow generalizability amongst all datasets however, it can be concurred that PCA works well for all datasets as it helps reduce dimensions more significantly and does not compromise a lot on variance capture which is the major goal of PCA itself. The major reason that these techniques work so well is because tabular data can be explained via linear mappings. Since PCA, LDA and SVD all are linear transformers, they can capture the hidden trends in the data well. However, same cannot be said for textual or image data where there are non-linear patterns.

6.1 Classification Datasets

Most of the techniques performed well but Linear discriminant analysis stood out for every dataset. LDA significantly reduced dimensions without compromising on variance capture making it an extremely useful technique. The major reason LDA can outperform the other techniques is because it is a supervised one which could be considered its drawback as well. However, since we are focused on classification, we would require a labeled dataset and LDA uses the labels along with the dataset to increase separability between classes. PCA only focuses on the linear mappings between the predictor variable whereas LDA focuses on linear mapping between the entire dataset and predicted variable as well. In addition, LDA's focus on class separability becomes the major factor that helps in improving classification since LDA focuses on capturing distinct information between the classes. However, LDA forcefully reduces dimensions to less than number of classes which may cause losing essential information if the entire dataset is relevant so we must resort to PCA or SVD if there is a major performance drop with LDA.

6.2 Regression Datasets

For regression datasets, LDA was not applicable as it requires a classification dataset. PCA worked well for these as it was able to capture most of the variance with significant feature reduction. Singular value decomposition performed like PCA as well where in some cases it was able to surpass PCA however PCA still remained the winner technique for regression datasets. The reason PCA works so well is because it can remove multicollinearity between variables and map the dataset on independent dimensions that capture the highest variance. Furthermore, because regression has a continuous output, it is more prone to having noise than classification datasets so amongst the other techniques PCA can handle noise better.

7. Conclusion

Overall, DR techniques must become a standard pre-processing step for high dimensional datasets to avoid unnecessary prolonged computation time and make machine learning simpler.